

ARTIFICIAL INTELLIGENCE¹

INTRODUCTION

¹The slides have been prepared using the textbook material available on the web, and the slides of the previous editions of the course by Prof. Luigia Carlucci Aiello

Summary

- ◇ What is AI?
- ◇ A short history
- ◇ The state of the art

What is AI?

“[The automation of] activities that we associate with human thinking, activities such as decision-making, problem-solving, learning ...” (Bellman 1978)

“The study of mental faculties through the use of computational models”
(Charniak McDermott 1985)

“The study of how to make computers do things at which, at the moment, people are better” (Rich Knight 1991)

“AI ... is concerned with intelligent behaviour in artifacts” (Nilsson, 1998)

What is AI?

Systems that think like humans	Systems that think rationally
Systems that act like humans	Systems that act rationally

Acting humanly: The Turing test

Turing (1950) “Computing machinery and intelligence”:

- ◇ “Can machines think?” → “Can machines behave intelligently?”
- ◇ Operational test for intelligent behavior: the Imitation Game
- ◇ Predicted that by 2000, a machine might have a 30% chance of fooling a lay person for 5 minutes
- ◇ Anticipated all major arguments against AI in following 50 years
- ◇ Suggested major components of AI: knowledge, reasoning, language understanding, learning

Problem: Turing test is not **reproducible**, **constructive**, or amenable to **mathematical analysis**

Thinking humanly: Cognitive Science

1960s “cognitive revolution”: information-processing psychology replaced prevailing orthodoxy of behaviorism

Requires scientific theories of internal activities of the brain

- What level of abstraction? “Knowledge” or “circuits”?

- How to validate? Requires

- 1) Predicting and testing behavior of human subjects (top-down)

- 2) Direct identification from neurological data (bottom-up)

Both approaches (roughly Cognitive Science and Cognitive Neuroscience) are now distinct.

Thinking rationally: Laws of Thought

Normative (or prescriptive) rather than descriptive

Aristotle: what are correct arguments/thought processes?

Several Greek schools developed various forms of logic:
notation and rules of derivation for thoughts;

Direct line through mathematics and philosophy to modern AI

Problems:

- 1) Formalize knowledge is difficult
- 2) How far can we get with computation of inferences?

Acting rationally

Rational behaviour: doing the right thing

The right thing: that which is expected to maximize goal achievement, given the available information

Does not necessarily involve thinking—e.g., blinking reflex—but thinking should be in the service of rational action

Aristotle (Nicomachean Ethics):

Every art and every inquiry, and similarly every action and pursuit, is thought to aim at some good

Rational agents

An **agent** is an entity that perceives and acts

This course is about designing **rational agents**

Abstractly, an agent is a function from percept histories to actions:

$$f : \mathcal{P}^* \rightarrow \mathcal{A}$$

For any given class of environments and tasks, we seek the agent (or class of agents) with the best performance

Caveat (1): *computational limitations make perfect rationality unachievable* \rightarrow design best **program** for given resources

Caveat (2): Herbert Simon' satisficing behaviour: applies to AI?

More mundane views

Intelligent = not stupid

where not is defined as negation as failure (to prove)

AI aims at solving difficult problems. When a problem is solved, it is no longer AI

AI is typically characterized a set of techniques.

AI and other disciplines

Philosophy	logic, methods of reasoning mind as physical system foundations of learning, language, rationality
Mathematics	formal representation and proof algorithms, computation, (un)decidability, (in)tractability, probability
Psychology	adaptation phenomena of perception and motor control experimental techniques (psychophysics, etc.)
Economics	formal theory of rational decisions
Linguistics	knowledge representation grammar
Neuroscience	plastic physical substrate for mental activity
Control theory	homeostatic systems, stability simple optimal agent designs

History of AI

1931

The Austrian Kurt Goedel shows that in first-order predicate logic all true statements are derivable. In higher order logics, on the other hand, there are true statements that are unprovable.

1937

Alan Turing points out the limits of intelligent machines with the halting problem.

1943

McCulloch and Pitts model neural networks and make the connection to propositional logic.

1950

Alan Turing defines machine intelligence with the Turing test and writes about learning machines and genetic algorithms

1951

Marvin Minsky develops a neural network machine. With 3000 vacuum tubes he simulates 40 neurons

History of AI

1955

Arthur Samuel (IBM) builds a learning chess program that plays better than its developer

1956

McCarthy organizes a conference in Dartmouth College. Here the name Artificial Intelligence was first introduced

Newell and Simon of Carnegie Mellon University (CMU) present the Logic Theorist, the first symbol-processing computer program

1958

McCarthy invents at the Massachusetts Institute of Technology (MIT) the high-level language LISP. He writes programs that are capable of modifying themselves

1959

Gelernter (IBM) builds the Geometry Theorem Prover

1961

The General Problem Solver (GPS) by Newell and Simon imitates human thought

History of AI

1963

McCarthy founds the AI Lab at Stanford University

1965

Robinson invents the resolution calculus for predicate logic

1966

Weizenbaum's program Eliza carries out dialogue with people in natural language

1969

Minsky and Papert show in their book Perceptrons that the perceptron, a very simple neural network, can only represent linear functions

1972

French scientist Alain Colmerauer invents the logic programming language PROLOG

British physician de Dombal develops an expert system for diagnosis of acute abdominal pain. It goes unnoticed in the mainstream AI community of the time

History of AI

1976

Shortliffe and Buchanan develop MYCIN, an expert system for diagnosis of infectious diseases, which is capable of dealing with uncertainty

1981

Japan begins, at great expense, the Fifth Generation Project with the goal of building a powerful PROLOG machine

1982

R1, the expert system for configuring computers, saves Digital Equipment Corporation 40 million dollars per year

1986

Renaissance of neural networks through, among others, Rumelhart, Hinton and Sejnowski. The system Nettetalk learns to read texts aloud

1990

Pearl, Cheeseman, Whittaker, Spiegelhalter bring probability theory into AI with Bayesian networks

Multi-agent systems become popular

The last 20 years of AI

1997 Deep Blue wins against Kasparov

2005 First DARPA Challenge on Autonomous CAR

2011 Watson wins Jeopardy champions

2016 Deepmind wins against GO champions

2017 Autonomous taxi

The state of the art: Competitions/Challenges

1997 First RoboCup in Nagoya

2018 22nd RoboCup in Montreal

- DARPA Grand Challenges
- European Robotics League
- Everywhere (ICAPS, ... Angry bird)

The state of the art: main areas

- Automated Reasoning
- Machine Learning
- Speech understanding and NLP
- Robotics
- Vision

The state of the art: keywords

AI and the Web (AIW) AI for Social Impact (AIS) Applications (APP) Cognitive Modeling (CM) Cognitive Systems (CS) Computational Sustainability (CSUS) Constraint Satisfaction and Optimization (CSO) Game Playing and Interactive Entertainment (GPIE) Game Theory and Economic Paradigms (GTEP) Heuristic Search and Optimization (HSO) Human-AI Collaboration (HAC) Humans and AI (HAI) Human-Computation and Crowd Sourcing (HCC) Knowledge Representation and Reasoning (KRR) Machine Learning (ML) Multiagent Systems (MAS) Natural Language Processing (NLP) Planning, Routing, and Scheduling (PRS) Reasoning under Uncertainty (RU) Robotics (ROB) Vision (VIS)

The state of the art: Applications

APP: Natural Sciences APP: Art and Music APP: Biomedical/Bioinformatics/Medicine APP: Computational Social Science APP: Computer-Aided Education APP: Economic/Financial APP: Energy APP: Entertainment APP: Environmental APP: Game Playing APP: Humanities APP: Intelligent User Interfaces APP: Marketing APP: Mobility APP: Networks APP: Philosophical and Ethical Issues APP: Security and Privacy APP: Social Networks APP: Transportation APP: Web APP: Other Applications

State of the art

Big data and Deep Learning made a breakthrough

- speech understanding
- image classification
- autonomous cars

Philosophical questions

- ◇ Can a machine **think** ?
- ◇ Can a machine have an **intelligent behaviour** ?
- ◇ Can a machine have **consciousness** ?
- ◇ Can a machine have **self-consciousness** ?
- ◇ Can a machine have **emotions, ...** ?

Ethical issues

- ◇ Advantages and disadvantages of **automation**
- ◇ Impact on **human intelligence**?
- ◇ **Privacy**
- ◇ **Legal Issues**
- ◇ **Ethics**

Asimov's laws of robotics

1. A robot may not injure a human being or, through inaction, allow a human being to come to harm;
2. a robot must obey the orders given it by human beings, except when such orders would conflict with the previous law;
3. a robot must protect its own existence as long as such protection does not conflict with the previous two laws.

Oren Etzioni's laws of AI

1. An A.I. system must be subject to the full gamut of laws that apply to its human operator.
2. An A.I. system must clearly disclose that it is not human.
3. An A.I. system cannot retain or disclose confidential information without explicit approval from the source of that information.

Clarke's laws

1. When a distinguished but elderly scientist states that something is possible, he is almost certainly right. When he states that something is impossible, he is very probably wrong.
2. The only way of discovering the limits of the possible is to venture a little way past them into the impossible.
3. Any sufficiently advanced technology is indistinguishable from magic.