

SEARCH IN THE STATE SPACE¹

LECTURE 3

¹The slides have been prepared using the textbook material available on the web, and the slides of the previous editions of the course by Prof. Luigia Carlucci Aiello

Problem solving

Summary(Russell&Norvig Chap. 3)

- ◇ Problem solving agents
- ◇ Types of problems
- ◇ Problem formulation
- ◇ Example problems
- ◇ Basic search algorithms

Agent Architectures

- **simple-reflex agents** (**stimoulus–response, tropis-tic**): react to stimuli
- **model-based agents** build a representation of the world
- **goal-based agents** have a goal and try to achieve it
- **utility-based agents** evaluate the utility of (goal) states according to a performance measure and trying to maximize it

Design of goal-based agents

- **search**: translating knowledge into a specialized representation of **states** and a set of **operators**
- **planning**: explicit representation of the state and specialized techniques for reasoning.
- **inference**: general state representation and reasoning.

The problem solving agent

- **problem formulation**:
 - state definition, the set of states is called **state space**
 - **initial state**
 - **operators** (actions characterizing state transitions); **path** in the state space is a sequence of operators
- **goal formulation**: set of states or goal test
- **search**: process to find a **solution**, i.e. a path from the initial state to one of the goal states.
- **execution**: the solution is performed by the agent

function SIMPLE-PS-AGENT(*percept*) **returns** an action

static: *seq*, an action sequence, initially empty

state, description of the current world state

goal, a goal, initially null

problem, a problem formulation

state \leftarrow UPDATE-STATE(*state*, *percept*)

if *seq* is empty **then**

goal \leftarrow FORMULATE-GOAL(*state*)

problem \leftarrow FORMULATE-PROBLEM(*state*, *goal*)

seq \leftarrow SEARCH(*problem*)

action \leftarrow RECOMMENDATION(*seq*, *state*)

seq \leftarrow REMAINDER(*seq*, *state*)

return *action*

Offline vs Online Problem Solving

The agent implements *offline* problem solving;
solution executed “eyes closed.”

Online problem solving involves acting without complete knowledge.

Example: Romania

On holiday in Romania; currently in Arad.
Flight leaves tomorrow from Bucharest

Formulate goal:

be in Bucharest

Formulate problem:

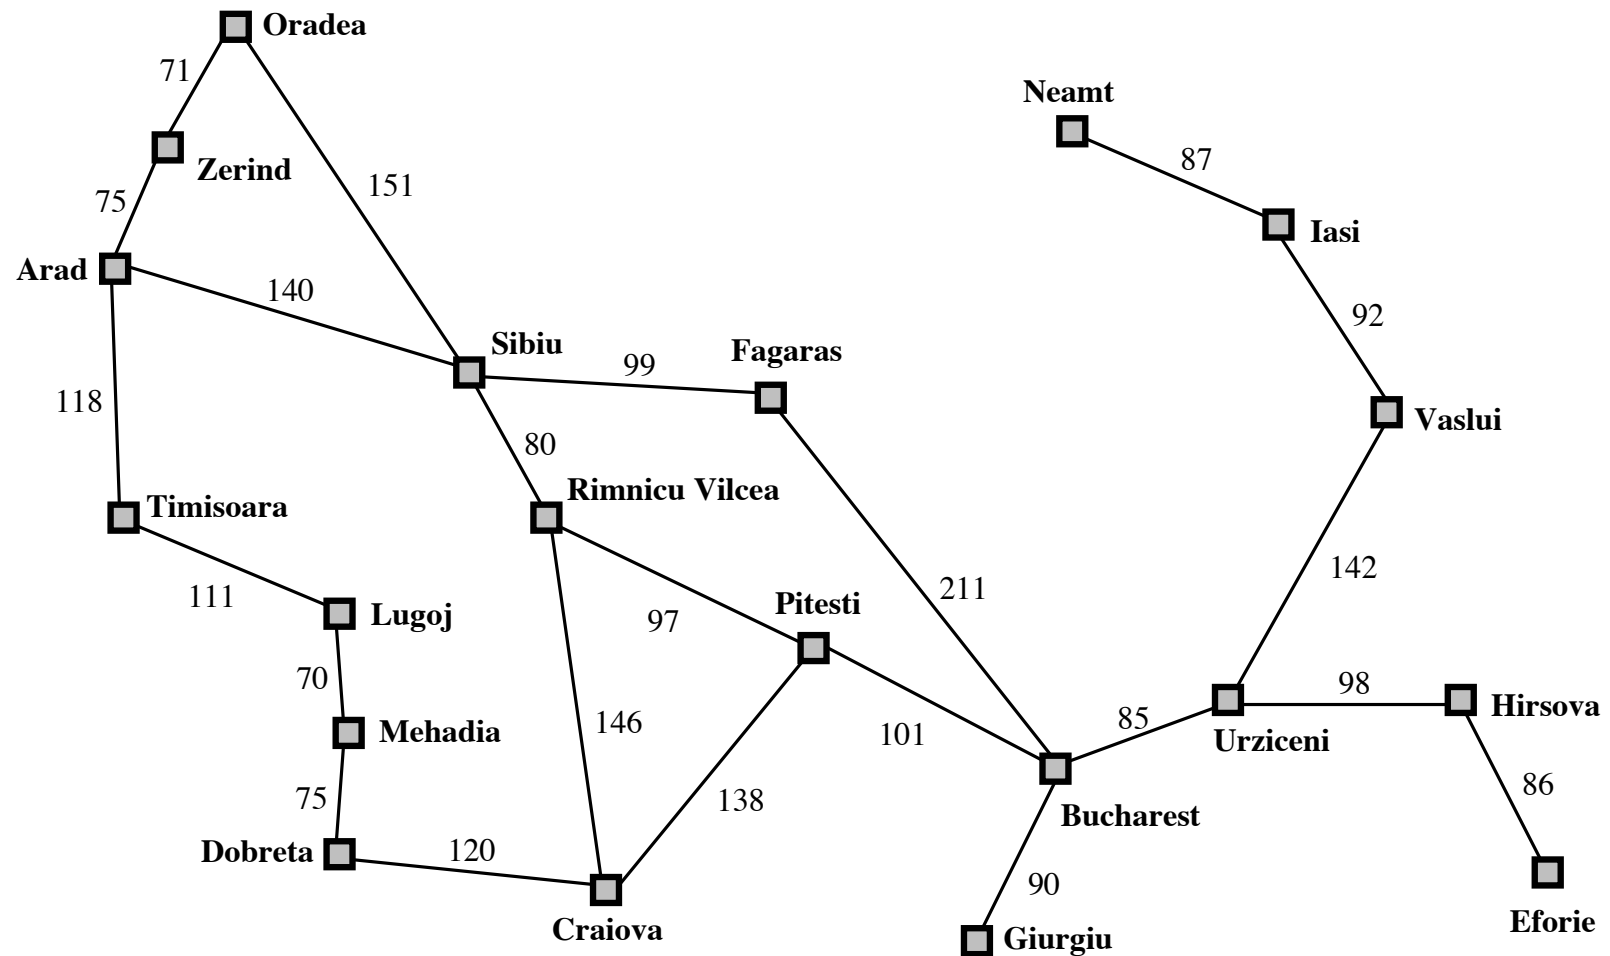
states: various cities

actions: drive between cities

Find solution:

sequence of cities, e.g., Arad, Sibiu, Fagaras, Bucharest

Example: Romania



Single-state problem formulation

A *problem* is defined by four items:

initial state e.g., “at Arad”

successor function $S(x)$ = set of action–state pairs
e.g., $S(Arad) = \{\langle Arad \rightarrow Zerind, Zerind \rangle, \dots\}$

goal test, can be

explicit, e.g., $x = \text{“at Bucharest”}$

implicit, e.g., $NoDirt(x)$

path cost (additive)

e.g., sum of distances, number of actions executed, etc.

$c(x, a, y)$ is the *step cost*, assumed to be ≥ 0

solution = sequence of actions from init state to goal state

Selecting a state space

Real world is absurdly complex

⇒ state space must be *abstracted* for problem solving

(Abstract) state = set of real states


(Abstract) action = complex combination of real actions

e.g., “Arad → Zerind” represents a complex set of possible routes, detours, rest stops, etc.

(Abstract) solution =

set of real paths that are solutions in the real world

Example: vacuum world state space graph



states: integer dirt and robot locations (ignore dirt *amounts*)

actions: *Left, Right, Suck, NoOp*

goal test: no dirt

path cost: 1 per action (0 for *NoOp*)

[

Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

states: locations of tiles (ignore intermediate positions)

actions: move blank left, right, up, down (ignore unjam etc.)

goal test: = goal state (given)


path cost: 1 per move

[Note: optimal solution of n -Puzzle family is NP-hard]

Real world problems

- Find a flight itinerary
- VLSI design
- Mobile robot motion
- Assembly sequences
- Internet search

Example: robotic assembly



states: real-valued coordinates of
robot joint angles
parts of the object to be assembled

actions: continuous motions of robot joints

goal test: complete assembly

path cost: time to execute

Tree search algorithms

Basic idea:

offline, simulated exploration of state space
by generating successors of already-explored states
(a.k.a. *expanding* states)

function TREE-SEARCH(*problem*, *strategy*) **returns** a solution, or failure

 initialize the search tree to the initial state of *problem*

loop do

if there are no candidates for expansion

then return failure

 choose a leaf node for expansion based on *strategy*

if the node contains a goal state


then return the corresponding solution

else expand the node and


 add the resulting nodes to the search tree

end

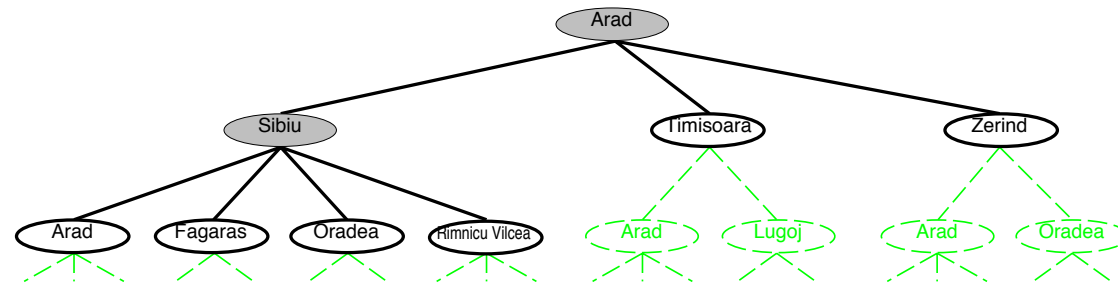
Tree search algorithms



Tree search algorithms




Tree search algorithms



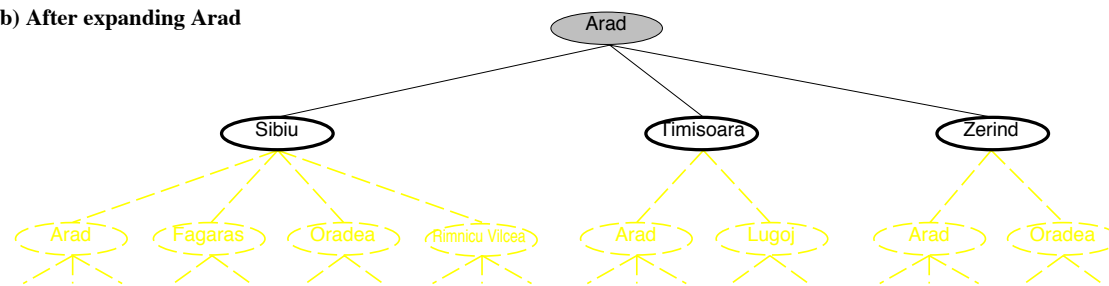
The search tree

- The **search tree** is composed by the search nodes and models the search process.
- **Expansion** is the process that builds successor nodes by applying the OPERATORS.
- The choice of the node to be expanded represents the **search strategy**.

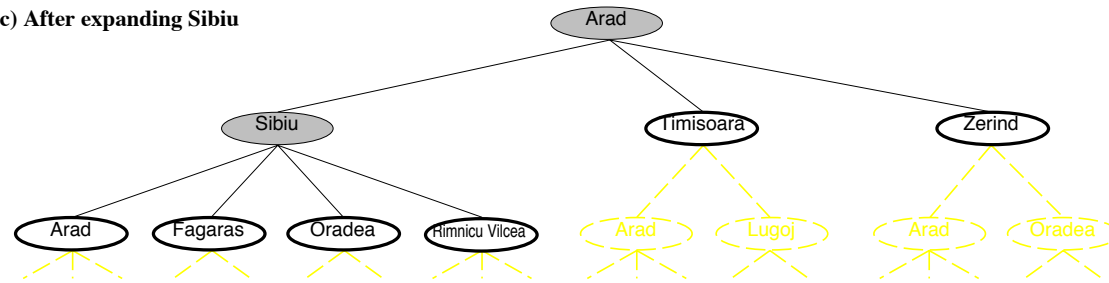
(a) The initial state



(b) After expanding Arad



(c) After expanding Sibiu




Implementation: states vs. nodes

A *state* is a (representation of) a physical configuration

A *node* is a data structure constituting part of a search tree

includes *parent*, *children*, *depth*, *path cost* $g(x)$

States do not have parents, children, depth, or path cost!



Implementation: general tree search

function TREE-SEARCH(*problem*, *fringe*) **returns** a solution, or failure

fringe \leftarrow INSERT(MK-NODE(INIT-STATE[*problem*]),
 fringe)

loop do

if *fringe* is empty **then return** failure

node \leftarrow REMOVE-FRONT(*fringe*)

if GOAL-TEST[*problem*] on STATE(*node*) **succeeds**
 then return *node*

fringe \leftarrow INSALL(EXPAND(*node*, *problem*), *fringe*)

function EXPAND(*node*, *problem*) **returns** a set of nodes

successors \leftarrow the empty set

for each *action*, *result*

in SUCCESSOR-FN[*problem*](STATE[*node*]) **do**

s \leftarrow a new NODE

PARENT-NODE[*s*] \leftarrow *node*;

ACTION[*s*] \leftarrow *action*;

STATE[*s*] \leftarrow *result*

PATH-COST[*s*] \leftarrow PATH-COST[*node*] +
STEP-COST(*node*, *action*, *s*)

DEPTH[*s*] \leftarrow DEPTH[*node*] + 1

add *s* to *successors*

return *successors*

Search strategies

A strategy is defined by picking the *order of node expansion*

Strategies are evaluated along the following dimensions:

completeness—does it always find a solution if one exists?

time complexity—number of nodes generated/expanded

space complexity—maximum number of nodes in memory

optimality—does it always find a least-cost solution?

Time and space complexity are measured in terms of

b —maximum branching factor of the search tree

d —depth of the least-cost solution

m —maximum depth of the state space (may be ∞)

Uninformed search strategies

Uninformed strategies use only the information available in the problem definition


- Breadth-first search
- Uniform-cost search
- Depth-first search
- Depth-limited search
- Iterative deepening search

Breadth-first search

Expand shallowest unexpanded node

Implementation:

fringe is a FIFO queue, i.e., new successors go at end




Breadth-first search

Expand shallowest unexpanded node

Implementation:

fringe is a FIFO queue, i.e., new successors go at end




Breadth-first search

Expand shallowest unexpanded node

Implementation:

fringe is a FIFO queue, i.e., new successors go at end




Breadth-first search

Expand shallowest unexpanded node

Implementation:

fringe is a FIFO queue, i.e., new successors go at end



Properties of breadth-first search

Complete Yes (if b is finite)

Time $1 + b + b^2 + b^3 + \dots + b^d + b(b^d - 1) = O(b^{d+1})$, i.e.,
exp. in d

Space $O(b^{d+1})$ (keeps every node in memory)

Optimal Yes (if cost = 1 per step); not optimal in general

Space is the big problem; can easily generate nodes at 10MB/sec
so 24hrs = 860GB.

Uniform-cost search

Expand least-cost unexpanded node

Implementation:

fringe = queue ordered by path cost

Equivalent to breadth-first if step costs all equal

Complete Yes, if step cost $\geq \epsilon$

Time # of nodes with $g \leq$ cost of optimal solution, $O(b^{\lceil (1+C^*/\epsilon) \rceil})$,
where C^* is the cost of the optimal solution

Space # of nodes with $g \leq$ cost of optimal solution, $O(b^{\lceil C^*/\epsilon \rceil})$


Optimal Yes—nodes expanded in increasing order of $g(n)$

Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front




Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front




Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front




Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front




Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front




Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front




Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front




Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front




Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front



Properties of depth-first search

Complete No: fails in infinite-depth spaces, spaces with loops

Modify to avoid repeated states along path

⇒ complete in finite spaces

Time $O(b^m)$: terrible if m is much larger than d , but if solutions are dense, may be much faster than breadth-first

Space $O(bm)$, i.e., linear space!

Optimal No

Depth-limited search

= depth-first search with depth limit l ,
i.e., nodes at depth l have no successors

Recursive implementation:

```
function DEPTH-LIMITED-SEARCH(problem, limit) returns soln/fail/cutoff  
  REC-DLS(MAKE-NODE(INI-STATE[problem]), problem, limit)
```

```
function REC-DLS(node, problem, limit) returns soln/fail/cutoff
```

```
  cutoff-occurred?  $\leftarrow$  false
```

```
  if GOAL-TEST[problem](STATE[node]) then return node
```

```
  else if DEPTH[node] = limit then return cutoff
```

```
  else for each successor in EXPAND(node, problem) do
```

```
    result  $\leftarrow$  RECURSIVE-DLS(successor, problem, limit)
```

```
    if result = cutoff then cutoff-occurred?  $\leftarrow$  true
```

```
    else if result  $\neq$  failure then return result
```

```
  if cutoff-occurred? then return cutoff else return failure
```

Iterative deepening search

```
function ITERATIVE-DEEPENING-SEARCH( problem)  
returns a solution  
  inputs: problem, a problem  
  for depth  $\leftarrow$  0 to  $\infty$  do  
    result  $\leftarrow$  DEPTH-LIMITED-SEARCH( problem,  
    depth)  
    if result  $\neq$  cutoff then return result  
  end
```

Iterative deepening search $l = 0$

Limit = 0




Iterative deepening search $l = 1$

Limit = 1




Iterative deepening search $l = 2$

Limit = 2



Iterative deepening search $l = 3$



Properties of iterative deepening search

Complete Yes

Time $(d + 1)b^0 + db^1 + (d - 1)b^2 + \dots + b^d = O(b^d)$

Space $O(bd)$

Optimal Yes, if step cost = 1

Can be modified to explore uniform-cost tree

Comparison for $b = 10$ and $d = 5$, solution at far right:

$$N(\text{IDS}) = 50 + 400 + 3,000 + 20,000 + 100,000 = 123,450$$

$$N(\text{BFS}) = 10 + 100 + 1,000 + 10,000 + 100,000 + 999,990 = 1,111,000$$

Search direction

Forward, or data driven

Backward, or goal driven

Bidirectional, or mixed

how do we choose?

1. Invertible operators?

2. “Structure” of the state space: branching factor, number of goal states, ecc.


Example of invertible operator:

move one step right \rightarrow move one step left

Non invertible operators: chess mate

Bidirectional search

If a problem allows for both forward and backward search:
We start from the initial state in forward and from the goal state in backward, trying to meet “in the middle”



Properties of bidirectional search

Time $O(b^{d/2})$

Space $O(b^{d/2})$

Completeness ed **Optimality** depend on the techniques chosen for the search

Problems:

- check the meet nodes $O(b^{d/2})$
- goal states implicitly defined

Summary of algorithms

Crit	Breadth- First	Uniform- Cost	Depth- First	Depth- Limited	Iter Deep	Bid (if app)
Time	b^d	b^d	b^m	b^l	b^d	$b^{d/2}$
Space	b^d	b^d	bm	bl	bd	$b^{d/2}$
Optim	Yes	Yes	No	No	Yes	Yes
Compl	Yes	Yes	No	Yes, if $l \geq d$	Yes	Yes

b is the branching factor;

d is the depth of the solution;

m is the maximum depth of the search tree;


l is the depth limit.

State repetition

- Avoid cyclic paths. The expansion step (or the set of operators) must avoid to generate successors that coincide with any ancestor.
- Avoid the generation of already generated states. Every state must be recorded in memory: space complexity $O(b^d)$ ($O(s)$, where s is the number of states). Hash tables.

Repeated states

Failure to detect repeated states can turn a linear problem into an exponential one!



function GRAPH-SEARCH(*problem*, *fringe*) **returns** a solution, or failure

closed \leftarrow an empty set

fringe \leftarrow INSERT(MAKE-NODE(INIT-STATE[*problem*]),
fringe)

loop do

if *fringe* is empty **then return** failure

node \leftarrow REMOVE-FRONT(*fringe*)

if GOAL-TEST[*problem*](STATE[*node*])

then return *node*

if STATE[*node*] is not in *closed* **then**

add STATE[*node*] to *closed*

fringe \leftarrow INSERTALL(EXPAND(*node*, *problem*),
fringe)

end

Summarizing

Problem formulation must lead to a search space that can be systematically explored by applying suitable operators

Abstraction is essential in problem formulation

There are several uninformed search strategies

Iterative deepening search uses linear space
and not much more time than other blind methods

Informed search allows to solve only small size problems

Missionaries and Cannibals

“Suppose you have a raw boat” Robot ?

3 cannibals and 3 missionaries must cross a river with a small boat that can hold 2 passengers at most. The number of missionaries must always be more or equal wrt the number of cannibals, otherwise they are eaten by the cannibals.

How can the missionaries cross the river without being eaten?