

# Artificial Intelligence

## 14. Planning Formalisms

How to Describe Problems, and What is a “Problem” Anyway?

Prof Sara Bernardini  
bernardini@diag.uniroma1.it  
www.sara-bernardini.com



SAPIENZA  
UNIVERSITÀ DI ROMA

Autumn Term

# Agenda

- 1 Introduction
- 2 Transition Systems
- 3 STRIPS Planning
- 4 Finite-Domain Representation (FDR) Planning
- 5 STRIPS vs. FDR
- 6 Extended Planning Frameworks [for Reference]
- 7 Conclusion

# Reminder: Planning = General Problem Solving

(some new problem)



describe problem in planning language  $\mapsto$  use off-the-shelf solver



(its solution)

- Any problem that can be formulated as a **planning problem**.
- Don't write the C++ code, just describe the problem!
- Don't maintain the C++ code, maintain the description!

# What is a Planning Problem?

Given a **planning task**:

- A description of the **initial state**.
- A description of the **goal condition**.
- A description of a set of **possible actions**.

→ Find a schedule of actions (a **plan**) that brings us from the initial state to a state in which the goal condition holds.

# Classical Planning

... makes **Simplifying Assumptions**:

- Initial situation unique and completely known, environment deterministic, static, discrete, single-agent.
- Actions executed one-by-one, plans are sequences.

**This is often not the case in practice!** Examples? Handling uncertainty (robot control), temporal/parallel execution (transportation), ...

**So why do we do this?**

- Clean framework to study planning problems. (Simplicity is a virtue!)
- Where most influential ideas were conceived.
- Successful applications using classical planning.
- We can successfully **compile** many extended paradigms into classical planning.

→ We focus entirely on classical planning in this course.

# Algorithmic Problems in Planning

## Satisficing Planning

**Input:** A planning task  $\Pi$ .

**Output:** A plan for  $\Pi$ , or **unsolvable** if no plan for  $\Pi$  exists.

## Optimal Planning

**Input:** A planning task  $\Pi$ .

**Output:** An optimal plan for  $\Pi$ , or **unsolvable** if no plan for  $\Pi$  exists.

→ The techniques successful for either one of these are almost disjoint!

→ Satisficing planning is *much* more effective in practice.

→ Programs solving these problems are called (optimal) **planners**, **planning systems**, or **planning tools**.

# Computational Complexity in Planning

**Why?** From this course's point of view, it's simply one technical tool we need.

→ To get a heuristic  $h$ , we map the planning problem into a simpler (abstract/relaxed) planning problem, from whose solution we compute  $h$ . To compute  $h$  efficiently, the “simpler” problem must be solvable in polynomial time.

**Definition (PlanEx and PlanOpt).** *PlanEx* is the problem of deciding, given a (STRIPS or FDR) planning task  $\Pi$ , *whether or not there exists a plan for  $\Pi$* . *PlanOpt* is the problem of deciding, given  $\Pi$  and  $B \in \mathbb{R}_0^+$ , *whether or not there exists a plan for  $\Pi$  whose cost is at most  $B$* .

→ PlanEx  $\approx$  satisficing planning, PlanOpt  $\approx$  optimal planning.

**Theorem (Planning is Hard).** *Each of PlanEx and PlanOpt is PSPACE-complete.*

**Proof.** See Chapter 13 and Bylander (1994), if interested.

# Reminder: NP and PSPACE

**Def Turing machine:** Works on a **tape** consisting of **tape cells**, across which its **R/W head** moves. The machine has **internal states**. There are **transition rules** specifying, given the current cell content and internal state, what the subsequent internal state will be, how what the R/W head does (write a symbol and/or move). Some internal states are **accepting**.

**Def NP:** Decision problems for which there exists a *non-deterministic* Turing machine that runs in *time* polynomial in the size of its input. Accepts if *at least one* of the possible runs accepts.

**Def PSPACE:** Decision problems for which there exists a *deterministic* Turing machine that runs in *space* polynomial in the size of its input.

**Relation:** Non-deterministic polynomial space can be simulated in deterministic polynomial space. Thus **PSPACE** = **NPSPACE**, and hence (trivially) **NP**  $\subseteq$  **PSPACE**. It is commonly believed that **NP**  $\not\subseteq$  **PSPACE** (similar to **P**  $\subseteq$  **NP**).

→ For comprehensive details, please see a text book. A good one is [Garey and Johnson (1979)]. (On the first 3 pages, they explain why knowing about NP-hardness will help you talk to your future boss.)



# Our Agenda for This Chapter

- ② **Transition Systems:** The basic framework we'll be moving in; forms the basis for both STRIPS and FDR. (= state space)
- ③ **STRIPS Planning:** STRIPS is by far the most wide-spread planning formalism. It is also the simplest possible reasonably expressive planning formalism, and thus a canonical subject to study.
- ④ **Finite-Domain Representations (FDR):** FDR is only slightly more general than STRIPS, but as we shall see can be quite useful.
- ⑤ **STRIPS vs. FDR:** The two formalisms can be compiled into each other. Such compilations are wide-spread in practice, and we will use them at some points during the course.
- ⑥ **Extended Planning Frameworks:** To at least give you a brief glimpse beyond classical planning.

# Transition Systems

→ State space of planning task = a transition system.

**Definition (Transition System).** A *transition system* is a 6-tuple

$\Theta = (S, L, c, T, I, S^G)$  where:

- $S$  is a finite set of *states*.
- $L$  is a finite set of transition *labels*.
- $c : L \mapsto \mathbb{R}_0^+$  is the *cost function*.
- $T \subseteq S \times L \times S$  is the *transition relation*.
- $I \in S$  is the *initial state*.
- $S^G \subseteq S$  is the set of *goal states*.

The *size* of  $\Theta$  is its number of states,  $\text{size}(\Theta) := |S|$ .

We say that  $\Theta$  *has the transition*  $(s, l, s')$  if  $(s, l, s') \in T$ . We also write this  $s \xrightarrow{l} s'$ , or  $s \rightarrow s'$  when not interested in  $l$ .

We say that  $\Theta$  is *deterministic* if, for all states  $s$  and labels  $l$ , there is at most one state  $s'$  with  $s \xrightarrow{l} s'$ .

We say that  $\Theta$  has *unit costs* if, for all  $l \in L$ ,  $c(l) = 1$ .

# Transition Systems, ctd.

**Terminology:**  $\Theta = (S, A, c, T, I, S^G)$ ;  $s, s', s_i \in S$


- $s'$  **successor** of  $s$  if  $s \rightarrow s'$ ;  $s$  **predecessor** of  $s'$  if  $s \rightarrow s'$ .
- $s'$  **reachable** from  $s$  if there exists a sequence of transitions:

$$s = s_0 \xrightarrow{l_1} s_1, \dots, s_{n-1} \xrightarrow{l_n} s_n = s'$$

- $n = 0$  possible; then  $s = s'$ .
- $l_1, \dots, l_n$  is called **path** from  $s$  to  $s'$ .
- $s_0, \dots, s_n$  is also called **path** from  $s$  to  $s'$ .
- The **cost** of that path is  $\sum_{i=1}^n c(l_i)$ .
- $s'$  **reachable** (without reference state) means reachable from  $I$ .
- **Solution** for  $s$ : path from  $s$  to some  $s' \in S^G$ ; **optimal** if cost is minimal among all solutions for  $s$ .
- $s$  is **solvable** if it has a solution; else,  $s$  is a **dead end**.
- Solution for  $I$  is called **solution for  $\Theta$** ;  $\Theta$  is **solvable** if it has a solution.

# Transition Systems: Illustration


Directed labeled graphs + mark-up for initial state and goal states:



- Are all states in  $\Theta$  reachable? No:  $s_7$
- Are all states in  $\Theta$  solvable? No:  $s_2$
- Is this  $\Theta$  deterministic? No: On two of the goal states ( $s_8, s_{10}$ ), actG labels more than one outgoing transition.


# Transition Systems: Illustration, ctd.

Directed labeled graphs + mark-up for initial state and goal states:



- **What are the optimal solutions for  $\Theta$ ?** Any path that starts with actB, applies actE  $n \in \{0, 2, 4, \dots\}$  times, then applies actC then actE and then no action other than actG.

# Why don't we simply use Dijkstra? Example Blocksworld



- $n$  blocks, 1 hand.
- A single action either takes a block with the hand or puts a block we're holding onto some other block/the table.

blocks	states	blocks	states
1	1	10	58941091
2	3	11	824073141
3	13	12	12470162233
4	73	13	202976401213
5	501	14	3535017524403
6	4051	15	65573803186921
7	37633	16	1290434218669921
8	394353	17	26846616451246353
9	4596553	18	588633468315403843

→ We are interested in solving **huge** transition systems, represented in a **compact** way as **planning tasks** (up next).

# STRIPS Planning: Syntax

**Definition (STRIPS Planning Task).** A *STRIPS planning task* is a 5-tuple  $\Pi = (P, A, c, I, G)$  where:

- $P$  is a finite set of *facts*, also *propositions*.
- $A$  is a finite set of *actions*; each  $a \in A$  is a triple  $a = (pre_a, add_a, del_a)$  of subsets of  $P$  referred to as the action's *precondition*, *add list*, and *delete list* respectively; we require that  $add_a \cap del_a = \emptyset$ .
- $c : A \mapsto \mathbb{R}_0^+$  is the *cost function*.
- $I \subseteq P$  is the *initial state*.
- $G \subseteq P$  is the *goal*.

We say that  $\Pi$  has *unit costs* if, for all  $a \in A$ ,  $c(a) = 1$ . We will often give each action  $a \in A$  a *name* (a string), and identify  $a$  with that name.

→ **Why do we allow 0-cost actions?** Negligible cost (e.g. switch light on, take photo with smartphone), asking questions about only one kind of actions (e.g. Mars rover *take-picture* only).

# STRIPS Encoding of “TSP” in Australia



- **Propositions**  $P$ :  $\{at(x), visited(x) \mid x \in \{Sydney, Adelaide, Brisbane, Perth, Darwin\}\}$ .
- **Initial state**  $I$ :  $\{at(Sydney), visited(Sydney)\}$ .
- **Goal**  $G$ :  $\{at(Sydney)\} \cup \{visited(x) \mid x \in \{Sydney, Adelaide, Brisbane, Perth, Darwin\}\}$ .
- **Actions**  $a \in A$ :  $drive(x, y)$  where  $x, y$  have a road.
  - Precondition**  $pre_a$ :  $\{at(x)\}$ .
  - Add list**  $add_a$ :  $\{at(y), visited(y)\}$ .
  - Delete list**  $del_a$ :  $\{at(x)\}$ .
- **Cost function**  $c$ :
 
$$c(drive(x, y)) = \begin{cases} 1 & \{x, y\} = \{Sydney, Brisbane\} \\ 1.5 & \{x, y\} = \{Sydney, Adelaide\} \\ 3.5 & \{x, y\} = \{Adelaide, Perth\} \\ 4 & \{x, y\} = \{Adelaide, Darwin\} \end{cases}$$



# STRIPS Planning: Semantics

**Definition (STRIPS State Space).** Let  $\Pi = (P, A, c, I, G)$  be a STRIPS planning task. The *state space* of  $\Pi$  is the labeled transition system  $\Theta_\Pi = (S, L, c, T, I, S^G)$  where:

- The *states* (also *world states*)  $S = 2^P$  are the subsets of  $P$ .
- The labels  $L = A$  are  $\Pi$ 's actions; the cost function  $c$  is that of  $\Pi$ .
- The transitions are  $T = \{s \xrightarrow{a} s' \mid a \in A[s], s' = s[a]\}$ , where  
 $A[s] := \{a \in A \mid pre_a \subseteq s\}$  are the actions *applicable* in  $s$ ; for  $a \in A[s]$ ,  
 $s[a] := (s \setminus del_a) \cup add_a$ ; for  $a \notin A[s]$ ,  $s[a]$  is *undefined*,  $s[a] := \perp$ .
- The initial state  $I$  is identical to that of  $\Pi$ .
- The goal states  $S^G = \{s \in S \mid G \subseteq s\}$  are those that satisfy  $\Pi$ 's goal.

An (optimal) *plan* for  $s \in S$  is an (optimal) solution for  $s$  in  $\Theta_\Pi$ . A solution for  $I$  is called a *plan for  $\Pi$* .  $\Pi$  is *solvable* if a plan for  $\Pi$  exists.

For  $\vec{a} = \langle a_1, \dots, a_n \rangle$ ,  $s[\vec{a}] := \begin{cases} s & n = 0 \\ s[\langle a_1, \dots, a_{n-1} \rangle][a_n] & n > 0 \end{cases}$

→ Is  $\Theta_\Pi$  *deterministic*? Yes: the successor state  $s'$  in  $s \xrightarrow{a} s'$  is uniquely determined by  $s$  and  $a$ , through  $s' = s[a]$ .

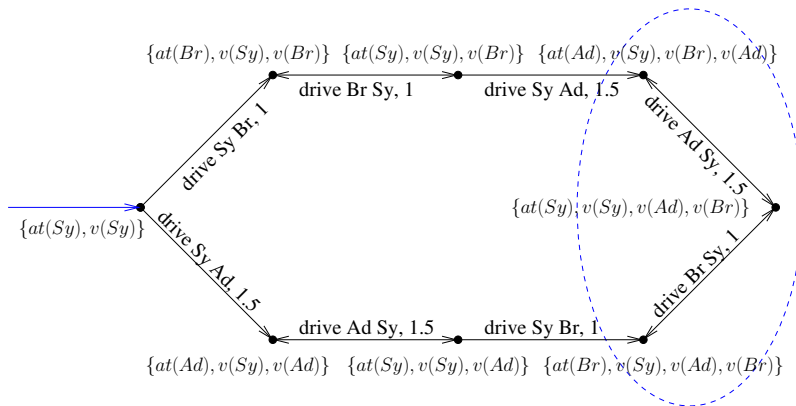
# STRIPS Encoding of Simplified “TSP”



- **Propositions**  $P$ :  $\{at(x), visited(x) \mid x \in \{Sydney, Adelaide, Brisbane\}\}$ .
- **Initial state**  $I$ :  $\{at(Sydney), visited(Sydney)\}$ .
- **Goal**  $G$ :  $\{visited(x) \mid x \in \{Sydney, Adelaide, Brisbane\}\}$ . (Note: no “ $at(Sydney)$ ”.)
- **Actions**  $a \in A$ :  $drive(x, y)$  where  $x, y$  have a road.
  - Precondition**  $pre_a$ :  $\{at(x)\}$ .
  - Add list**  $add_a$ :  $\{at(y), visited(y)\}$ .
  - Delete list**  $del_a$ :  $\{at(x)\}$ .
- **Cost function**  $c$ :

$$c(drive(x, y)) = \begin{cases} 1 & \{x, y\} = \{Sydney, Brisbane\} \\ 1.5 & \{x, y\} = \{Sydney, Adelaide\} \end{cases}$$

# STRIPS Encoding of Simplified “TSP”: State Space



→ Exactly one optimal plan: drive Sy Br, drive Br Sy, drive Sy Ad.

→ **Is this actually the state space?** No, only the reachable part. E.g.,  $\Theta_{\Pi}$  also includes the states  $\{v(Sy)\}$  and  $\{at(Sy), at(Br)\}$ .

# Questionnaire



- **Propositions**  $P$ :  
 $\{at(x), visited(x) \mid x \in \{Sydney, Adelaide, Brisbane, Perth, Darwin\}\}.$
- **Initial state**  $I$ :  $\{at(Sydney), visited(Sydney)\}.$

How many states are there in the “TSP in Australia” task?

→:  $2^{10} = 1024$ . But only a small portion of them are reachable (less than  $5 \cdot 2^4 = 80$ )!

# FDR Planning: Syntax

**Definition (FDR Planning Task).** A *finite-domain representation planning task*, short *FDR planning task*, is a 5-tuple  $\Pi = (V, A, c, I, G)$  where:

- $V$  is a finite set of *state variables*, each  $v \in V$  with a finite domain  $D_v$ . We refer to (partial) functions on  $V$ , mapping each  $v \in V$  into a member of  $D_v$ , as (partial) *variable assignments*.
- $A$  is a finite set of *actions*; each  $a \in A$  is a pair  $(pre_a, eff_a)$  of *partial variable assignments* referred to as the action's *precondition* and *effects*.
- $c : A \mapsto \mathbb{R}_0^+$  is the *cost function*.
- $I$  is a complete variable assignment called the *initial state*.
- $G$  is a partial variable assignment called the *goal*.

We say that  $\Pi$  has *unit costs* if, for all  $a \in A$ ,  $c(a) = 1$ .

→ In FDR, a (partial) variable assignment represents a state in  $I$ , a condition in  $pre_a$  and  $G$ , and an effect instruction in  $eff_a$ .

**Notation:** Pairs  $(v, d)$  are *facts*, also written  $v = d$ . We identify partial variable assignments  $p$  with fact sets. We write  $V[p] := \{v \in V \mid p(v) \text{ is defined}\}$ .

# FDR Encoding of “TSP”



- **Variables**  $V$ :  $at : \{Sydney, Adelaide, Brisbane, Perth, Darwin\}$ ;  $visited(x) : \{T, F\}$  for  $x \in \{Sydney, Adelaide, Brisbane, Perth, Darwin\}$ .
- **Initial state**  $I$ :  $at = Sydney$ ,  $visited(Sydney) = T$ ,  $visited(x) = F$  for  $x \neq Sydney$ .
- **Goal**  $G$ :  $at = Sydney$ ,  $visited(x) = T$  for all  $x$ .
- **Actions**  $a \in A$ :  $drive(x, y)$  where  $x, y$  have a road.  
     **Precondition**  $pre_a$ :  $at = x$ .  
     **Effect**  $eff_a$ :  $at = y$ ,  $visited(y) = T$ .

- **Cost function**  $c$ :

$$c(drive(x, y)) = \begin{cases} 1 & \{x, y\} = \{Sydney, Brisbane\} \\ 1.5 & \{x, y\} = \{Sydney, Adelaide\} \\ 3.5 & \{x, y\} = \{Adelaide, Perth\} \\ 4 & \{x, y\} = \{Adelaide, Darwin\} \end{cases}$$

# FDR Planning: Semantics

**Definition (FDR State Space).** Let  $\Pi = (V, A, c, I, G)$  be an FDR planning task. The *state space* of  $\Pi$  is the labeled transition system

$\Theta_\Pi = (S, L, c, T, I, S^G)$  where:

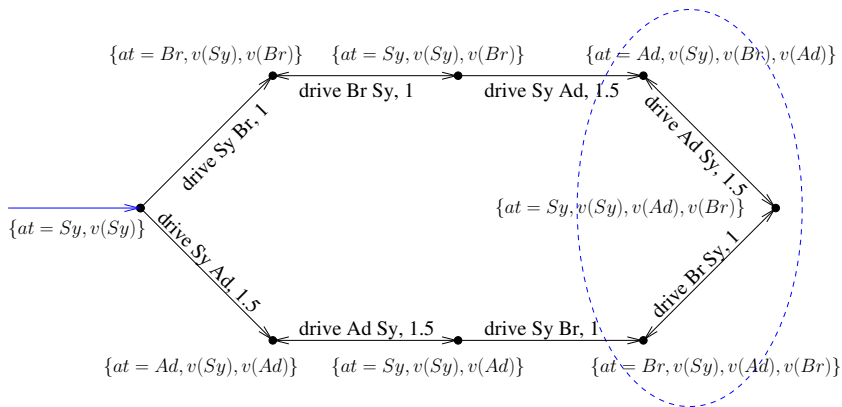
- The *states* (also *world states*)  $S$  are the complete variable assignments.
- The labels  $L = A$  are  $\Pi$ 's actions; the cost function  $c$  is that of  $\Pi$ .
- The transitions are  $T = \{s \xrightarrow{a} s' \mid a \in A[s], s' = s[a]\}$ , where  
 $A[s] := \{a \in A \mid pre_a \subseteq s\}$  are the actions *applicable* in  $s$ ; for  $a \notin A[s]$ ,  
 $s[a] := \perp$ ; for  $a \in A[s]$ ,  $s[a](v) := \begin{cases} eff_a(v) & v \in V[eff_a] \\ s(v) & v \notin V[eff_a] \end{cases}$
- The initial state  $I$  is identical to that of  $\Pi$ .
- The goal states  $S^G = \{s \in S \mid G \subseteq s\}$  are those that satisfy  $\Pi$ 's goal.

→ In  $s[a]$ , instead of “adding/deleting” facts, we overwrite the previous variable values by  $eff_a$ .

→ Plan, optimal plan,  $s[\vec{a}]$  for action sequence  $\vec{a}$ : as before (slide 20).

# FDR Encoding of Simplified “TSP”: State Space

(using “ $v(x)$ ” as shorthand for  $visited(x) = T$ )



→ This is only the reachable part of the state space: E.g.,  $\Theta_{\Pi}$  also includes the state  $\{at = Sy, v(Br)\}$ . (But neither  $\{v(Sy)\}$  nor  $\{at = Sy, at = Br\}$ , compare slide 22.)



# Questionnaire

## Question!

**How many STRIPS state variables are needed to encode the problem of finding a path in a graph with  $n$  vertices?**

(A): 1

(B):  $n$

(C):  $\lceil \log_2 n \rceil$

(D):  $2 * \lceil \log_2 n \rceil$

→ (D): We need to encode our current position in the graph. This can be done with  $n$  propositions of the form "*at(p)*", but it can be done more compactly by: numbering the positions  $ID(p)$ ; representing  $ID(p)$  in the binary system using  $\lceil \log_2 n \rceil$  bits  $bit_i$ ; and representing each  $bit_i$  with two STRIPS facts  $True(bit_i)$  and  $False(bit_i)$ .

## Question!

**How many FDR state variables are needed for this?**

(A): 1

(B):  $n$

(C):  $\lceil \log_2 n \rceil$

(D):  $2 * \lceil \log_2 n \rceil$

→ (A): We need 1 variable with  $n$  values, encoding our current position in the graph.

# STRIPS vs. FDR in Practice

## How do people use FDR?

- Our surface language is PDDL, which corresponds to STRIPS.
- Most implemented planning tools are based on **Fast Downward (FD)** [Helmert (2009)], which reads PDDL input, then internally uses a “clever” **STRIPS-2-FDR translation** (see next).
- That translation involves a **PSPACE**-complete sub-problem.

## Why??? Practical Efficiency!

- **Regression**: FDR avoids myriads of unreachable states. → **Chapter 17**
- **Causal Graphs**: Capture variable dependencies; have a much clearer structure for clever FDR (e.g., acyclic vs. cyclic). → **Chapter 16**
- **Complexity Analysis**: Better with clearer causal graph. → **Chapter 16**
- **Construction of Heuristic Functions**: Better with multiple-valued variables and clearer causal graph. → **Chapters 18**
- **Modeling**: Anyway, FDR is more natural! (It's just one truck, after all.)

## Why does anybody use STRIPS? It's a legacy system.

→ We should be modeling in FDR. For historical reasons, we aren't.

# STRIPS vs. FDR Conversions

## Conversions:

- ❶ **FDR-2-STRIPS:** For each variable  $v$  with domain  $\{d_1, \dots, d_k\}$ , make  $k$  STRIPS facts " $v = d_1$ ",  $\dots$ , " $v = d_k$ ".
- ❷ **STRIPS-2-FDR:** Naïve vs. clever variants, see slides 36 – 39.

## What role does all this play here?

- Both STRIPS and FDR are used in practice. The **programming exercises** focus on the planner Fast Downward, which uses FDR.
- Some techniques in the remainder of the course are easier to introduce in STRIPS, some are easier in FDR, so we will keep both around.
- Specific **relevance of (I):** If the course introduces a technique  $A$  in STRIPS, then  $A$  in FDR (and hence your FD code!) is equivalent to "convert-FDR-2-STRIPS-then-do- $A$ ".
- Specific **relevance of (II):** So you get an understanding of how FD processes the PDDL/STRIPS input to FDR.

# Isomorphism

**Definition (Isomorphism).** Let  $\Theta = (S, L, c, T, I, S^G)$  and  $\Theta' = (S', L', c', T', I', S'^G)$  be transition systems. We say that  $\Theta$  is *isomorphic* to  $\Theta'$ , written  $\Theta \sim \Theta'$ , if there exist bijective functions  $\varphi : S \mapsto S'$  and  $\psi : L \mapsto L'$  such that:

- (i)  $\varphi(I) = I'$ .
- (ii)  $s \in S^G$  iff  $\varphi(s) \in S'^G$ .
- (iii)  $(s, l, t) \in T$  iff  $(\varphi(s), \psi(l), \varphi(t)) \in T'$ .
- (iv) For all  $l \in L$ ,  $c(l) = c'(\psi(l))$ .

→ Isomorphic transition systems are identical modulo renaming states and actions.

→ Isomorphisms typically result from compilations between different formalisms (see later this chapter); we will also sometimes use them as a technical device.

# FDR-2-STRIPS: Details

**Definition (FDR-2-STRIPS).** Let  $\Pi = (V, A, c, I, G)$  be an FDR planning task. The *STRIPS conversion* of  $\Pi$  is the STRIPS task

$\Pi^{\text{STR}} = (P_V, A^{\text{STR}}, c, I, G)$  where:

- $P_V = \{v = d \mid v \in V, d \in D_v\}$  is the set of (STRIPS) facts.
- $A^{\text{STR}} = \{a^{\text{STR}} \mid a \in A\}$  where  $\text{pre}_{a^{\text{STR}}} = \text{pre}_a$ ,  $\text{add}_{a^{\text{STR}}} = \text{eff}_a$ , and  

$$\text{del}_{a^{\text{STR}}} = \bigcup_{(v=d) \in \text{eff}_a} \begin{cases} \{v = \text{pre}_a(v)\} & \text{if } \text{pre}_a(v) \text{ is defined} \\ \{v = d' \mid d' \in D_v \setminus \{d\}\} & \text{otherwise} \end{cases}$$
- The cost function  $c$  is defined by  $c(a^{\text{STR}}) := c(a)$  for all  $a^{\text{STR}} \in A^{\text{STR}}$ .
- $I$  and  $G$  are identical to those of  $\Pi$ .

→ The adds establish the new variable values of  $\text{eff}_a$ ; the deletes make sure to erase the previous values of those variables.

→ Take-home message: **FDR variable/value pairs  $\approx$  STRIPS facts!**

**Proposition.** Let  $\Pi = (V, A, c, I, G)$  be an FDR planning task, and let  $\Pi^{\text{STR}}$  be its STRIPS conversion. Then  $\Theta_\Pi$  is isomorphic to the sub-system of  $\Theta_{\Pi^{\text{STR}}}$  induced by those  $s \subseteq P_V$  where, for each  $v \in V$ ,  $s$  contains exactly one fact of the form  $v = d$ . All other states in  $\Theta_{\Pi^{\text{STR}}}$  are unreachable.

# FDR-2-STRIPS: Simplified “TSP”



- **FDR  $V$ :**  $at : \{Sydney, Adelaide, Brisbane\}; visited(x) : \{T, F\}$  for  $x \in \{Sydney, Adelaide, Brisbane\}$ .
- **STRIPS  $P$ :**  $at(x), visited(x, T), visited(x, F)$  for  $x \in \{Sydney, Adelaide, Brisbane\}$ .
- **FDR  $dr(x, y)$ :**  $pre = \{at = x\}, eff = \{at = y, v(y) = T\}$ .
- **STRIPS  $dr(x, y)$ :**  
 $pre = \{at(x)\}, add = \{at(y), v(y, T)\}, del = \{at(x), v(y, F)\}$ .

# STRIPS-2-FDR: Naïve Translation

**Definition (STRIPS-2-FDR).** Let  $\Pi = (P, A, c, I, G)$  be a STRIPS planning task. The *FDR conversion* of  $\Pi$  is the FDR task

$\Pi^{\text{FDR}} = (V_P, A^{\text{FDR}}, c, I^{\text{FDR}}, G^{\text{FDR}})$  where:

- $V_P = \{v_p \mid p \in P\}$  is the set of variables, *all Boolean*.
- $A^{\text{FDR}} = \{a^{\text{FDR}} \mid a \in A\}$  where  $\text{pre}_{a^{\text{FDR}}} = \{v_p = T \mid p \in \text{pre}_a\}$  and  $\text{eff}_{a^{\text{FDR}}} = \{v_p = T \mid p \in \text{add}_a\} \cup \{v_p = F \mid p \in \text{del}_a\}$ .
- The cost function  $c$  is defined by  $c(a^{\text{FDR}}) := c(a)$  for all  $a^{\text{FDR}} \in A^{\text{STR}}$ .
- $I = \{v_p = T \mid p \in I\}$ ; and  $G = \{v_p = T \mid p \in G\}$ .

→ All variables here have two possible values only, so this does not benefit at all from the added expressivity of FDR. Hence the designation “naïve”.

**Proposition.** Let  $\Pi = (P, A, c, I, G)$  be a STRIPS planning task, and let  $\Pi^{\text{FDR}}$  be its STRIPS conversion. Then  $\Theta_\Pi$  is isomorphic to  $\Theta_{\Pi^{\text{STR}}}$ .

# STRIPS-2-FDR, Naïve: Simplified “TSP”



- STRIPS  $P$ :  $at(x)$ ,  $visited(x)$  for  $x \in \{Sydney, Adelaide, Brisbane\}$ .
- FDR  $V$ :  $at(x)$ ,  $visited(x) : \{T, F\}$  for  $x \in \{Sydney, Adelaide, Brisbane\}$ .
- STRIPS  $dr(x, y)$ :  $pre = \{at(x)\}$ ,  $add = \{at(y), v(y)\}$ ,  $del = \{at(x)\}$
- FDR  $dr(x, y)$ :  $pre = \{at(x) = T\}$ ,  
 $eff = \{at(y) = T, v(y) = T, at(x) = F\}$ .



# STRIPS-2-FDR: Clever Translation

## How to be clever?

- Find sets  $\{p_1, \dots, p_k\}$  of STRIPS facts so that every reachable state  $s$  makes exactly one  $p_i$  true.  
→ Deciding whether this holds, for a given  $\{p_1, \dots, p_k\}$ , is **PSPACE**-complete (cf. slide 31). But one can design fast algorithms finding *some* such sets [Helmert (2009)].
- For each set  $\{p_1, \dots, p_k\}$  found, make *one* FDR variable  $v$  with domain  $\{d_1, \dots, d_k\}$ .
- This is implemented in the pre-processor of **Fast Downward**.

# STRIPS-2-FDR Naïve vs. Clever: Simplified “TSP”



- **STRIPS**  $P$ :  $at(x), visited(x)$  for  $x \in \{Sydney, Adelaide, Brisbane\}$ .
- **Naïve**  $V$ :  $at(x), visited(x) : \{T, F\}$  for  $x \in \{Sydney, Adelaide, Brisbane\}$ .
- **Clever**  $V$ :  $at : \{Sydney, Adelaide, Brisbane\}$ ;  
 $visited(x) : \{T, F\}$  for  $x \in \{Sydney, Adelaide, Brisbane\}$ .

→ The naïve version is merely STRIPS in disguise. The clever version is more natural, and is explicit about the “truck position”.

# Action Description Language (ADL)

**Framework Definition:** [Pednault (1989); Hoffmann and Nebel (2001)].

**Problem:** Like STRIPS but with **first-order logic** (FOL) formulas in  $pre_a$  and  $G$ , and **conditional effects** that execute only if their individual effect condition holds.

**Plan:** Sequence of actions. (Yes, this is still “classical planning”.)

**Example:** If your action  $a$  opens the doors of an elevator, then each passenger gets out iff their individual condition (“Is this my destination floor?”) holds. If you want to satisfy complex constraints (“Group A should never meet group B in the elevator”) then  $pre_a$  gets nasty. (See the file miconic-ADL on Moodle.)

**Compilation:** FOL formulas: Ground them (the universe is finite) and transform to DNF [Gazen and Knoblock (1997); Koehler and Hoffmann (2000)].

Conditional effects: Either enumerate all combinations of effects, or introduce artificial facts/actions enforcing an “effect evaluation phase” [Nebel (2000)].

**State of the art:** Get rid of FOL formulas but keep the conditional effects [Hoffmann and Nebel (2001)].

# Numeric and Temporal Planning



**Numeric Planning:** [Fox and Long (2003)]

$pre_a : fuelSupply \geq distance(x, y) * fuelConsumption$

$eff_a : fuelSupply := fuelSupply - distance(x, y) * fuelConsumption$

**Compilation:** Nothing known.

**Temporal Planning:** [Fox and Long (2003)]

$duration_a : distance(x, y) / speed$

$eff_a : at\ Start \neg at(x), at\ End\ at(y).$

**Compilation:** Ignore durations during search, schedule plan as a post-process [Edelkamp (2003)]. **Competitive with state of the art!**

# Soft Goals and Trajectory Constraints



**Soft Goals:** [Gerevini *et al.* (2009)]

“I don’t absolutely have to visit Darwin, but if I do, I get a certain amount  $R$  of reward.”

**Compilation:** Artificial actions that allow to forgo each weak goal, at cost  $R$ ; minimize cost [Keyder and Geffner (2009)]. **State of the art!**

**Trajectory Constraints:** [Gerevini *et al.* (2009)]

“I must visit Perth before I visit Darwin.”

**Compilation:** Artificial preconditions/effects, e.g. *visited(Perth)* into precondition of driving to Darwin [Edelkamp (2006)]. **State of the art!**

# Conformant Planning

**Framework Definition:** [Smith and Weld (1998); Bonet and Givan (2006)].

**Problem:** There are **many possible initial states** (represented as a formula), and each action may have **several possible effects**. We have **no observability** during plan execution.

**Plan:** Sequence of actions that **achieves the goal regardless which initial state and action effects occur**.

**Example:** You're in a dark cave but don't know where exactly. The plan is to walk to the right until you reach a wall and can locate yourself. Then navigate to your goal by counting your steps.

**Compilation:** Artificial "what-if" facts, like "If I was at A initially, then I am now at B" [Palacios and Geffner (2009)]. **State of the art!**

# Contingent Planning

**Framework Definition:** e.g., [Hoffmann and Brafman (2005)].

**Problem:** There are many possible initial states (represented as a formula), and each action may have several possible effects. We have **partial observability** during plan execution.

**Plan:** **Tree of actions** that achieves the goal in each of its leaves. (“Plan ahead for all possible contingencies, i.e., situation aspects not known at planning time.”)

**Example:** Solving the Wumpus world: You walk some steps, then use sensing (for breeze and stench), and continue depending on the outcome.

**Compilation:** Sample initial states, classical planning with artificial facts encoding knowledge yields a plan tree for those; in case a problem is detected during execution, re-plan with the new state of knowledge [Shani and Brafman (2011)]. **Competitive with state of the art!**

# Probabilistic Planning

**Framework Definition:** e.g., [Younes *et al.* (2005)].

**Problem:** Each action specifies a **probability distribution over its possible effects**. We have **full observability** during plan execution. (**Markov Decision Process (MDP)** framework.)

**Plan:** **Policy** that maps states to actions in a way that maximizes the expected reward.

**Example:** Controlling a robot: If navigation comes with an imprecision (which it usually does), then the outcome of a “move” operation is uncertain.

**Compilation:** Make classical problem that acts as if you could *choose* the outcomes; find a plan, and execute; if the plan fails, then re-plan from the current state [Yoon *et al.* (2007)]. **State of the art for problems where “reactive behavior” is suitable** (things may go wrong, but if they do, they can be easily repaired).



# Summary

- **Transition systems** are a kind of directed graph (typically huge) that encode how the state of the world can change.
- **Planning tasks** are compact representations for transition systems, based on state variables; they are the input for **planning systems**.
- In **satisficing planning**, we must find a solution to planning tasks (or show that no solution exists). In **optimal planning**, we must additionally guarantee that generated solutions are the cheapest possible.
- **Classical planning** makes strong simplifying assumptions, but is very successful in practice and can be used by **compilation** to tackle more expressive planning problems.
- In **STRIPS**, state variables are Boolean; in **FDR**, they may have arbitrary finite domains. The two formalisms can be compiled into each other. FDR is preferable, but current planning technology is based on STRIPS for historical reasons.  
→ **PDDL**, see **Next Chapter**.

# Remarks

## Regarding the name “FDR”:

- FDR is not consistently named in the literature.
- It is often referred to as **SAS<sup>+</sup>** because that's what some complexity guys called it, in the first papers considering a formalism equivalent to our FDR [e.g., Bäckström and Nebel (1995)].
- [Helmert (2006)] called it **multi-valued planning tasks (MPT)** which can still be seen in some papers.
- [Helmert (2009)] finally called it FDR.

# Reading

- *Concise Finite-Domain Representations for PDDL Planning Tasks* [Helmert (2009)].

[Available on Moodle](#)

**Content:** Describes in detail the “clever” STRIPS-2-FDR conversion implemented in Fast Downward. The sets  $\{p_1, \dots, p_k\}$  of STRIPS facts, of which exactly one is true in every reachable state, are found by automatic [invariance analysis](#). Is in wide-spread use, and a basic familiarity with it is relevant for anybody working in planning.

# References I

Christer Bäckström and Bernhard Nebel. Complexity results for SAS<sup>+</sup> planning. *Computational Intelligence*, 11(4):625–655, 1995.

Blai Bonet and Robert Givan. 5th international planning competition: Non-deterministic track – call for participation. In *Proceedings of the 5th International Planning Competition (IPC'06)*, 2006.

Tom Bylander. The computational complexity of propositional STRIPS planning. *Artificial Intelligence*, 69(1–2):165–204, 1994.

Stefan Edelkamp. Taming numbers and durations in the model checking integrated planning system. *Journal of Artificial Intelligence Research*, 20:195–238, 2003.

Stefan Edelkamp. On the compilation of plan constraints and preferences. In Derek Long and Stephen Smith, editors, *Proceedings of the 16th International Conference on Automated Planning and Scheduling (ICAPS'06)*, pages 374–377, Ambleside, UK, 2006. AAAI Press.

Maria Fox and Derek Long. PDDL2.1: An extension to PDDL for expressing temporal planning domains. *Journal of Artificial Intelligence Research*, 20:61–124, 2003.

Michael R. Garey and David S. Johnson. *Computers and Intractability—A Guide to the Theory of NP-Completeness*. Freeman, San Francisco, CA, 1979.

# References II

B. Cenk Gazen and Craig Knoblock. Combining the expressiveness of UCPOP with the efficiency of Graphplan. In S. Steel and R. Alami, editors, *Proceedings of the 4th European Conference on Planning (ECP'97)*, pages 221–233. Springer-Verlag, 1997.

Alfonso Gerevini, Patrik Haslum, Derek Long, Alessandro Saetti, and Yannis Dimopoulos. Deterministic planning in the fifth international planning competition: PDDL3 and experimental evaluation of the planners. *Artificial Intelligence*, 173(5-6):619–668, 2009.

Malte Helmert. The Fast Downward planning system. *Journal of Artificial Intelligence Research*, 26:191–246, 2006.

Malte Helmert. Concise finite-domain representations for PDDL planning tasks. *Artificial Intelligence*, 173:503–535, 2009.

Jörg Hoffmann and Ronen Brafman. Contingent planning via heuristic forward search with implicit belief states. In Susanne Biundo, Karen Myers, and Kanna Rajan, editors, *Proceedings of the 15th International Conference on Automated Planning and Scheduling (ICAPS-05)*, pages 71–80, Monterey, CA, USA, 2005. AAAI Press.

## References III

- Jörg Hoffmann and Bernhard Nebel. The FF planning system: Fast plan generation through heuristic search. *Journal of Artificial Intelligence Research*, 14:253–302, 2001.
- Emil Keyder and Hector Geffner. Soft goals can be compiled away. *Journal of Artificial Intelligence Research*, 36:547–556, 2009.
- Jana Koehler and Jörg Hoffmann. On the instantiation of ADL operators involving arbitrary first-order formulas. In *Proceedings ECAI-00 Workshop on New Results in Planning, Scheduling and Design*, 2000.
- Bernhard Nebel. On the compilability and expressive power of propositional planning formalisms. *Journal of Artificial Intelligence Research*, 12:271–315, 2000.
- Hector Palacios and Hector Geffner. Compiling uncertainty away in conformant planning problems with bounded width. *Journal of Artificial Intelligence Research*, 35:623–675, 2009.
- Edwin P.D. Pednault. ADL: Exploring the middle ground between STRIPS and the situation calculus. In R. Brachman, H. J. Levesque, and R. Reiter, editors, *Principles of Knowledge Representation and Reasoning: Proceedings of the 1st International Conference (KR-89)*, pages 324–331, Toronto, ON, May 1989. Morgan Kaufmann.

# References IV

- Guy Shani and Ronen I. Brafman. Replanning in domains with partial information and sensing actions. In Toby Walsh, editor, *Proceedings of the 22nd International Joint Conference on Artificial Intelligence (IJCAI'11)*, pages 2021–2026. AAAI Press/IJCAI, 2011.
- D. E. Smith and D. Weld. Conformant Graphplan. In Jack Mostow and Charles Rich, editors, *Proceedings of the 15th National Conference of the American Association for Artificial Intelligence (AAAI'98)*, pages 889–896, Madison, WI, USA, July 1998. MIT Press.
- Sung Wook Yoon, Alan Fern, and Robert Givan. FF-Replan: a baseline for probabilistic planning. In Mark Boddy, Maria Fox, and Sylvie Thiebaux, editors, *Proceedings of the 17th International Conference on Automated Planning and Scheduling (ICAPS'07)*, pages 352–359, Providence, Rhode Island, USA, 2007. Morgan Kaufmann.
- Håkan L. S. Younes, Michael L. Littman, David Weissman, and John Asmuth. The first probabilistic track of the international planning competition. *Journal of Artificial Intelligence Research*, 24:851–887, 2005.