

Artificial Intelligence

13. Introduction to Planning

How to Describe Arbitrary Search Problems

Prof Sara Bernardini

bernardini@diag.uniroma1.it

www.sara-bernardini.com




SAPIENZA
UNIVERSITÀ DI ROMA

Spring Term

Agenda

- 1 Introduction
- 2 The History of Planning
- 3 The STRIPS Planning Formalism
- 4 Why Complexity Analysis?
- 5 Planning Complexity
- 6 Conclusion

Reminder: Classical Search Problems (Chapters 3 & 4)



- **States:** Card positions ($position_Jspades = Qhearts$).
- **Actions:** Card moves ($move_Jspades_Qhearts_freecell4$).
- **Initial state:** Start configuration.
- **Goal states:** All cards “home”.
- **Solution:** Card moves solving this game.

Planning

Ambition:

Write one program that can solve all classical search problems.

Types of description

- The **blackbox description** of a problem Π is an API (a programming interface) providing functionality allowing to construct the state space: `InitialState()`, `GoalTest(s)`, ...
→ "Specifying the problem" = programming the API.
 - The **declarative description** of Π comes in a **problem description language**. This allows to implement the API, and much more.
→ "Specifying the problem" = writing a problem description.
- Here, "problem description language" = **planning language**.

“Planning Language”?

How does a planning language describe a problem?

- A *logical description* of the possible **states** (vs. Blackbox: data structures). E.g.: predicate $Eq(.,.)$.
- A *logical description* of the **initial state** I (vs. data structures). E.g.: $Eq(x, 1)$.
- A *logical description* of the **goal condition** G (vs. a goal-test function). E.g.: $Eq(x, 2)$.
- A *logical description* of the set A of **actions** in terms of **preconditions** and **effects** (vs. functions returning applicable actions and successor states).
E.g.: “increment x : pre $Eq(x, 1)$, eff $Eq(x, 2) \wedge \neg Eq(x, 1)$ ”.

→ Solution (**plan**) = sequence of actions from A , transforming I into a state that satisfies G . E.g.: “increment x ”.


“Planning Language”?

Disclaimer:

→ Planning languages go way beyond classical search problems. There are variants for inaccessible, stochastic, dynamic, continuous, and multi-agent settings.

- We focus on classical search for simplicity (combined with practical relevance).
- For a comprehensive overview, see [Ghallab *et al.* (2004)].

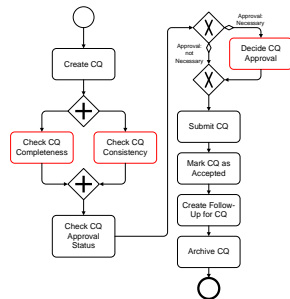
Application: Natural Language Generation



- **Input:** Tree-adjoining grammar, intended meaning.
- **Output:** Sentence expressing that meaning.


Application: Business Process Templates at SAP

Action name	precondition	effect
Check CQ Completeness	CQ.archiving:notArchived	CQ.completeness:complete OR CQ.completeness:notComplete
Check CQ Consistency	CQ.archiving:notArchived	CQ.consistency:consistent OR CQ.consistency:notConsistent
Check CQ Approval Status	CQ.archiving:notArchived AND CQ.approval:notChecked AND CQ.completeness:complete AND CQ.consistency:consistent	CQ.approval:necessary OR CQ.approval:notNecessary
Decide CQ Approval	CQ.archiving:notArchived AND CQ.approval:necessary	CQ.approval:granted OR CQ.approval:notGranted
Submit CQ	CQ.archiving:notArchived AND (CQ.approval:notNecessary OR CQ.approval:granted)	CQ.submission:submitted
Mark CQ as Accepted	CQ.archiving:notArchived AND CQ.submission:submitted	CQ.acceptance:accepted
Create Follow-Up for CQ	CQ.archiving:notArchived AND CQ.acceptance:accepted	CQ.followUp:documentCreated
Archive CQ	CQ.archiving:notArchived	CQ.archiving:archived




- **Input:** SAP-scale model of behavior of activities on Business Objects, process endpoint.
- **Output:** Process template leading to this point.

Application: Automatic Hacking




- **Input:** Network configuration, location of sensible data.
- **Output:** Sequence of exploits giving access to that data.




Planning!



Action name	precondition	effect
Check CQ Completeness	CQ.archiving:notArchived	CQ.completeness:complete OR CQ.completeness:notComplete
Check CQ Consistency	CQ.archiving:notArchived	CQ.consistency:consistent OR CQ.consistency:notConsistent
Check CQ Approval Status	CQ.archiving:notArchived AND CQ.approval:notChecked AND CQ.completeness:complete AND CQ.consistency:consistent	CQ.approval:necessary OR CQ.approval:notNecessary
Decide CQ Approval	CQ.archiving:notArchived AND CQ.approval:necessary	CQ.approval:granted OR CQ.approval:notGranted
Submit CQ	CQ.archiving:notArchived AND (CQ.approval:notNecessary OR CQ.approval:granted)	CQ.submission:submitted
Mark CQ as Accepted	CQ.archiving:notArchived AND CQ.submission:submitted	CQ.acceptance:accepted
Create Follow-Up for CQ	CQ.archiving:notArchived AND CQ.acceptance:accepted	CQ.followUp:documentCreated
Archive CQ	CQ.archiving:notArchived	CQ.archiving:archived



Planning Domain Definition Language (PDDL) \mapsto Planning System




Reminder: General Problem Solving, Pros and Cons

- **Powerful:** In some applications, generality is absolutely necessary. (E.g. SAP)
- **Quick:** Rapid prototyping: 10s lines of problem description vs. 1000s lines of C++ code. (E.g. language generation)
- **Flexible:** Adapt/maintain *the description*. (E.g. network security)
- **Intelligent:** Determines automatically how to solve a complex problem effectively! (The ultimate goal, no?!)
- **Efficiency loss:** Without any domain-specific knowledge about Chess, you don't beat Kasparov ...
→ Trade-off between “automatic and general” vs. “manualwork but effective”.

How to make fully automatic algorithms effective?

ps. “Making Fully Automatic Algorithms Effective”



- n blocks, 1 hand.
- A single action either takes a block with the hand or puts a block we're holding onto some other block/the table.

blocks	states	blocks	states
1	1	9	4596553
2	3	10	58941091
3	13	11	824073141
4	73	12	12470162233
5	501	13	202976401213
6	4051	14	3535017524403
7	37633	15	65573803186921
8	394353	16	1290434218669921

→ State spaces typically are huge even for simple problems.

→ In other words: Even solving “simple problems” automatically (without help from a human) requires a form of intelligence. **With blind search, even the largest super-computer in the world won't scale beyond 20 blocks!**

Algorithmic Problems in Planning

Satisficing Planning

Input: A planning task Π .

Output: A plan for Π , or “unsolvable” if no plan for Π exists.

Optimal Planning

Input: A planning task Π .

Output: An *optimal* plan for Π , or “unsolvable” if no plan for Π exists.

→ The techniques successful for either one of these are almost disjoint.
And satisficing planning is *much* more effective in practice.

→ Programs solving these problems are called (optimal) **planners**, **planning systems**, or **planning tools**.

In the Beginning ...

... Man invented Robots:

"Planning" as in "the making of plans by an autonomous robot".

In a little more detail:

- Newell and Simon (1963) introduced **general problem solving**.
- ... *not much happened (well not much we still speak of today)* ...
- Stanford Research Institute developed a robot named **"Shakey"**.
- They needed a **"planning"** component taking decisions.
- They took inspiration from general problem solving and theorem proving, and called the resulting algorithm **"STRIPS"** (Stanford Research Institute Problem Solver).

And then:

History of Planning Algorithms

Compilation into Logics/Theorem Proving:

- **Popular when:** Stone Age – 1990.
- **Approach:** *From planning task description, generate FOL formula φ that is satisfiable iff there exists a plan; use a theorem prover on φ .*
- **Keywords/cites:** Situation calculus, frame problem, ...

Partial-Order Planning:

- **Popular when:** 1990 – 1995.
- **Approach:** *Starting at [goal](#), extend partially ordered set of actions by inserting achievers for open sub-goals, or by adding ordering constraints to avoid conflicts.*
- **Keywords/cites:** UCPOP [Penberthy and Weld (1992)], [causal links](#), [flaw-selection strategies](#), ...

History of Planning Algorithms, ctd.

GraphPlan:

- **Popular when:** 1995 – 2000.
- **Approach:** *In a forward phase, build a layered “planning graph” whose “time steps” capture which pairs of actions can achieve which pairs of facts; in a backward phase, search this graph starting at goals and excluding options proved to not be feasible.*
- **Keywords/cites:** [Blum and Furst (1995, 1997); Koehler et al. (1997)], [action/fact mutexes](#), [step-optimal plans](#), ...

Planning as SAT:

- **Popular when:** 1996 – today.
- **Approach:** *From planning task description, generate propositional CNF formula φ_k that is satisfiable iff there exists a plan with k steps; use a SAT solver on φ_k , for different values of k .*
- **Keywords/cites:** [Kautz and Selman (1992, 1996); Rintanen et al. (2006); Rintanen (2010)], [SAT encoding schemes](#), [BlackBox](#), ...

History of Planning Algorithms, ctd.

Planning as Heuristic Search:

- **Popular when:** 1999 – today.
- **Approach:** Devise a method \mathcal{R} to simplify (“relax”) any planning task Π ; given Π , solve $\mathcal{R}(\Pi)$ to generate a heuristic function h for informed search.
- **Keywords/cites:** [Bonet and Geffner (1999); Haslum and Geffner (2000); Bonet and Geffner (2001); Hoffmann and Nebel (2001); Edelkamp (2001); Gerevini *et al.* (2003); Helmert (2006); Helmert *et al.* (2007); Helmert and Geffner (2008); Karpas and Domshlak (2009); Helmert and Domshlak (2009); Richter and Westphal (2010); Nissim *et al.* (2011); Katz *et al.* (2012); Keyder *et al.* (2012); Katz *et al.* (2013); Domshlak *et al.* (2015)], **critical path heuristics, ignoring delete lists, relaxed plans, landmark heuristics, abstractions, partial delete relaxation, ...**

The International Planning Competition (IPC)

Competition?

“Run competing planners on a set of benchmarks devised by the IPC organizers. Give awards to the most effective planners.”

- 1998, 2000, 2002, 2004, 2006, 2008, 2011, 2014, 2018, 2023
- **PDDL** [McDermott *et al.* (1998); Fox and Long (2003); Hoffmann and Edelkamp (2005); Gerevini *et al.* (2009)]
- \approx 50 **domains**, \gg 1000 **instances**, 74 (!) planners in 2011
- **Optimal** track vs. **satisficing** track
- Various others: uncertainty, learning, ...

<http://ipc.icaps-conference.org/>

Winners of the IPC

Edition	Satisficing track	Optimal track
IPC'00:	FF, heuristic search	–
IPC'02:	LPG, heuristic search	–
IPC'04:	SGPlan, heuristic search	SATPLAN, SAT compilation
IPC'06:	SGPlan, heuristic search	SATPLAN, SAT compilation
IPC'08:	LAMA, heuristic search	Gamer, symbolic search
IPC'11:	LAMA, heuristic search	Fast-Downward, heuristic search
IPC'14:	IBACoP, portfolio	SymBA*, symbolic search
IPC'18:	FDSS, portfolio (HS)	Delfi portfolio (symbolic + HS)
IPC'23:	Scorpion-Maidu, portfolio (HS) Levitron, portfolio (HS)	Ragnarok portfolio (HS)

→ For the rest of this chapter, we focus on planning as heuristic search.

→ This is a **VERY** short summary of the history of the IPC! There are many different categories, and many different awards.

Questionnaire

Question!

If planners x, y both compete in IPC'YY, and x wins, is x “better than” y ?

(A): Yes.

(B): No.

→ Yes, but only on the IPC'YY benchmarks, and only according to the criteria used for determining a “winner”! On other domains and/or according to other criteria, you may well be better off with the “loser”.

→ It's complicated, over-simplification is dangerous. (But, of course, nevertheless is being done all the time).

“STRIPS” Planning

- **STRIPS** = Stanford Research Institute Problem Solver.
STRIPS is the simplest possible (reasonably expressive) logics-based planning language.
- STRIPS has only **Boolean variables**: propositional logic atoms.
- Its preconditions/effects/goals are as canonical as imaginable:
 - Preconditions, goals: **conjunctions** of **positive atoms**.
 - Effects: **conjunctions** of **literals** (positive or negated atoms).
- We use the common special-case notation for this simple formalism.
- I'll outline some extensions beyond STRIPS later on, when we discuss PDDL.

→ **Historical note:** STRIPS [Fikes and Nilsson (1971)] was originally a planner (cf. Shakey), whose language actually wasn't quite that simple.

STRIPS Planning: Syntax

Definition (STRIPS Planning Task). A *STRIPS planning task*, short *planning task*, is a 4-tuple $\Pi = (P, A, I, G)$ where:

- P is a finite set of *facts* (aka *propositions*).
- A is a finite set of *actions*; each $a \in A$ is a triple $a = (pre_a, add_a, del_a)$ of subsets of P referred to as the action's *precondition*, *add list*, and *delete list* respectively; we require that $add_a \cap del_a = \emptyset$.
- $I \subseteq P$ is the *initial state*.
- $G \subseteq P$ is the *goal*.

We will often give each action $a \in A$ a *name* (a string), and identify a with that name.

Note: We assume, for simplicity, that every action has cost 1. (**Unit costs**, cf. **Chapter 3**.)

“TSP” in Australia



[→ Strictly speaking, this is not actually a **TSP** problem instance; simplified/adapted for illustration.]

STRIPS Encoding of "TSP"



- **Facts** P : $\{at(x), visited(x) \mid x \in \{Sydney, Adelaide, Brisbane, Perth, Darwin\}\}$.
- **Initial state** I : $\{at(Sydney), visited(Sydney)\}$.
- **Goal** G :
 $\{at(Sydney)\} \cup \{visited(x) \mid x \in \{Sydney, Adelaide, Brisbane, Perth, Darwin\}\}$.
- **Actions** $a \in A$: $drive(x, y)$ where x, y have a road.
Precondition pre_a : $\{at(x)\}$.
Add list add_a : $\{at(y), visited(y)\}$.
Delete list del_a : $\{at(x)\}$.
- **Plan**: $\langle drive(Sydney, Brisbane), drive(Brisbane, Sydney), drive(Sydney, Adelaide), drive(Adelaide, Perth), drive(Perth, Adelaide), drive(Adelaide, Darwin), drive(Darwin, Adelaide), drive(Adelaide, Sydney) \rangle$.

STRIPS Planning: Semantics

Definition (STRIPS State Space). Let $\Pi = (P, A, c, I, G)$ be a STRIPS planning task. The *state space* of Π is $\Theta_\Pi = (S, A, T, I, S^G)$ where:


- The states (also *world states*) $S = 2^P$ are the subsets of P .
- A is Π 's action set.
- The transitions are $T = \{s \xrightarrow{a} s' \mid \text{pre}_a \subseteq s, s' = ([s], a)\}$.
If $\text{pre}_a \subseteq s$, then a is *applicable* in s and $([s], a) := (s \cup \text{add}_a) \setminus \text{del}_a$. If $\text{pre}_a \not\subseteq s$, then $([s], a)$ is undefined.
- I is Π 's initial state.
- The goal states $S^G = \{s \in S \mid G \subseteq s\}$ are those that satisfy Π 's goal.

An (optimal) *plan* for $s \in S$ is an (optimal) solution for s in Θ_Π , i.e., a path from s to some $s' \in S^G$. A solution for I is called a *plan for Π* . Π is *solvable* if a plan for Π exists.

For $\vec{a} = \langle a_1, \dots, a_n \rangle$, $([s], \vec{a}) := ([\dots](\big([s], a_1), a_2) \dots, a_n)$ if each a_i is applicable in the respective state; else, $([s], \vec{a})$ is undefined.

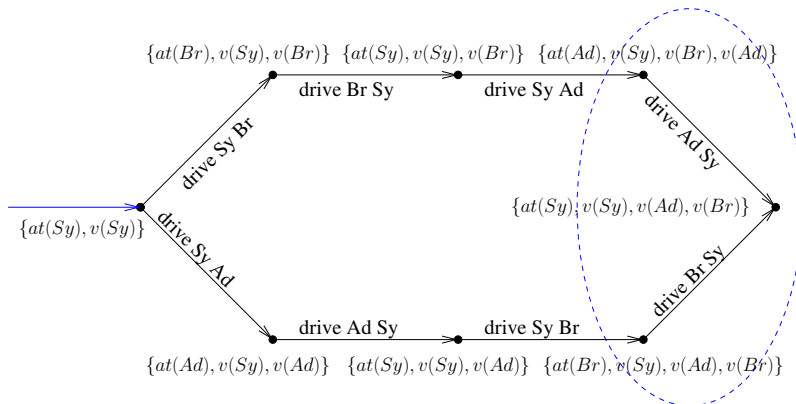
Note: This is exactly like the state spaces of **Chapter 3**, without a cost function. Solutions are defined as before (paths from I to a state in S^G).

STRIPS Encoding of Simplified “TSP”




- **Facts** P : $\{at(x), visited(x) \mid x \in \{Sydney, Adelaide, Brisbane\}\}$.
- **Initial state** I : $\{at(Sydney), visited(Sydney)\}$.
- **Goal** G : $\{visited(x) \mid x \in \{Sydney, Adelaide, Brisbane\}\}$. (Note: no “ $at(Sydney)$ ”.)
- **Actions** $a \in A$: $drive(x, y)$ where x, y have a road.
 - Precondition** pre_a : $\{at(x)\}$.
 - Add list** add_a : $\{at(y), visited(y)\}$.
 - Delete list** del_a : $\{at(x)\}$.

STRIPS Encoding of Simplified "TSP": State Space



→ Is this actually the state space? No, only the reachable part. E.g., Θ_{Π} also includes the states $\{v(Sy)\}$ and $\{at(Sy), at(Br)\}$.

(Oh no it's) The Blockworld



- **Facts:** $on(x, y)$, $onTable(x)$, $clear(x)$, $holding(x)$, $armEmpty()$.
- **Initial state:** $\{onTable(E), clear(E), \dots, onTable(C), on(D, C), clear(D), armEmpty()\}$.
- **Goal:** $\{on(E, C), on(C, A), on(B, D)\}$.
- **Actions:** $stack(x, y)$, $unstack(x, y)$, $putdown(x)$, $pickup(x)$.
- **$stack(x, y)$?** $pre : \{holding(x), clear(y)\}$
 $add : \{on(x, y), armEmpty()\}$
 $del : \{holding(x), clear(y)\}$.

Questionnaire

Question!

Which are correct encodings (ones that are part of some correct overall model) of the STRIPS Blocksworld *pickup(x)* action schema?

(A): (*onTable(x)*, *clear(x)*,
armEmpty()},
{*holding(x)*},
{*onTable(x)*}).

(C): (*onTable(x)*, *clear(x)*,
armEmpty()},
{*holding(x)*}, {*onTable(x)*,
armEmpty(), *clear(x)*}).

(B): (*onTable(x)*, *clear(x)*,
armEmpty()},
{*holding(x)*},
{*armEmpty()*}).

(D): (*onTable(x)*, *clear(x)*,
armEmpty()},
{*holding(x)*}, {*onTable(x)*,
armEmpty()}).

→ (A): No, must delete *armEmpty()*. (B): No, must delete *onTable(x)*. (C), (D): Both yes: We can, but don't have to, encode the *single-arm* Blocksworld so that the block currently in the hand is not clear. (For (C), *stack(x,y)* and *putdown(x)* need to add *clear(x)*, so the encoding on the previous slide does not work.)

Why Complexity Analysis?

Why? Why?

Two very good reasons:

- ① It saves you from spending lots of time trying to invent algorithms that do not exist.
- ② Killer app in planning: **tractable fragments for heuristic functions**.
 - Identify special cases that can be solved in polynomial time.
 - **Relax** the input into the special case to obtain a heuristic function!

→ I'll next remind you of the basic terms, then I'll illustrate both with an example. Afterwards we'll have a brief look at the complexity of the main decision problems in STRIPS planning.

Reminder (?): NP and PSPACE

Def Turing machine: Works on a **tape** consisting of **tape cells**, across which its **R/W head** moves. The machine has **internal states**. There are **transition rules** specifying, given the current cell content and internal state, what the subsequent internal state will be, how what the R/W head does (write a symbol and/or move). Some internal states are **accepting**.

Def NP: Decision problems for which there exists a *non-deterministic* Turing machine that runs in *time* polynomial in the size of its input. Accepts if *at least one* of the possible runs accepts.


Def PSPACE: Decision problems for which there exists a *deterministic* Turing machine that runs in *space* polynomial in the size of its input.

Relation: Non-deterministic polynomial space can be simulated in deterministic polynomial space. Thus **PSPACE** = **NPSPACE**, and hence (trivially) **NP** \subseteq **PSPACE**. It is commonly believed that **NP** $\not\subseteq$ **PSPACE** (similar to **P** \subseteq **NP**).

→ For comprehensive details, please see a text book. My personal favorite is [Garey and Johnson (1979)]. (On the first 3 pages, they explain why knowing about NP-hardness will help you talk to your future boss.)

The “Only-Adds” Relaxation

Example: “Logistics”



- **Facts P :** $\{truck(x) \mid x \in \{A, B, C, D\}\} \cup \{pack(x) \mid x \in \{A, B, C, D, T\}\}$.
- **Initial state I :** $\{truck(A), pack(C)\}$.
- **Goal G :** $\{truck(A), pack(D)\}$.
- **Actions A :** (Notated as “precondition \Rightarrow adds, \neg deletes”)
 - $drive(x, y)$, where x, y have a road:
“ $truck(x) \Rightarrow truck(y), \neg truck(x)$ ”.
 - $load(x)$: “ $truck(x), pack(x) \Rightarrow pack(T), \neg pack(x)$ ”.
 - $unload(x)$: “ $truck(x), pack(T) \Rightarrow pack(x), \neg pack(T)$ ”.

Only-Adds Relaxation: Drop the preconditions and deletes.

“ $drive(x, y): \Rightarrow truck(y)$ ”; “ $load(x): \Rightarrow pack(T)$ ”; “ $unload(x): \Rightarrow pack(x)$ ”.

→ Say we want to use this for generating a heuristic function.

Solving Only-Adds STRIPS Tasks

Our problem:

- Given STRIPS task $\Pi = (P, A, I, G)$.
- Find action sequence \vec{a} leading from I to a state that contains G , when **pretending that preconditions and deletes are empty**.

Solution 1: (simplest possible approach)

```
 $\vec{a} := \langle \rangle$   
while  $G \neq \emptyset$  do  
    select  $a \in A$   
     $G := G \setminus add_a$   
     $\vec{a} := \vec{a} \circ \langle a \rangle$ ;  $A := A \setminus \{a\}$   
endwhile  
return  $h := |\vec{a}|$ 
```

→ **Is this h admissible?** No. Admissibility is only guaranteed if we find a shortest possible \vec{a} ; else, \vec{a} might be longer than a plan for Π itself. Selecting an arbitrary action each time, \vec{a} may be longer than needed.

Solving Only-Adds STRIPS Tasks, ctd.

So, what about this?

```
 $\vec{a} := \langle \rangle$   
while  $G \neq \emptyset$  do  
  select  $a \in A$  s.t.  $|add_a|$  is maximal  
   $G := G \setminus add_a$   
   $\vec{a} := \vec{a} \circ \langle a \rangle$ ;  $A := A \setminus \{a\}$   
endwhile  
return  $h := |\vec{a}|$ 
```

→ h admissible? No, large add_a doesn't help if the intersection with G is small.

And this?

```
 $\vec{a} := \langle \rangle$   
while  $G \neq \emptyset$  do  
  select  $a \in A$  s.t.  $|add_a \cap G|$  is maximal  
   $G := G \setminus add_a$   
   $\vec{a} := \vec{a} \circ \langle a \rangle$ ;  $A := A \setminus \{a\}$   
endwhile  
return  $h := |\vec{a}|$ 
```

→ h admissible? Still no. Example: $G = \{A, B, C, D, E, F\}$; $add_{a_1} = \{A, B\}$;
 $add_{a_2} = \{C, D\}$; $add_{a_3} = \{E, F\}$; $add_{a_4} = \{A, C, E\}$.

Solving Only-Adds STRIPS Tasks, ctd.

From [Garey and Johnson (1979)]:

222

NP-COMPLETE PROBLEMS

[SP5] MINIMUM COVER

INSTANCE: Collection C of subsets of a finite set S , positive integer $K \leq |C|$.

QUESTION: Does C contain a cover for S of size K or less, i.e., a subset $C' \subseteq C$ with $|C'| \leq K$ such that every element of S belongs to at least one member of C' ?

Reference: [Karp, 1972]. Transformation from X3C.

Comment: Remains NP-complete even if all $c \in C$ have $|c| \leq 3$. Solvable in polynomial time by matching techniques if all $c \in C$ have $|c| \leq 2$.

So what?

- Given STRIPS task $\Pi = (P, A, I, G)$.
- Find **optimal** \vec{a} leading from I to a state that contains G , when pretending that preconditions and deletes are empty.

→ \vec{a} leads to $G \Leftrightarrow \bigcup_{a \in \vec{a}} add_a \supseteq G \Leftrightarrow$ the add lists in \vec{a} cover G . QED.

Questionnaire

Assume: In 3 years from now, you have finished your studies and are working in your first industry job. Your boss Mr. X gives you a problem and says “Solve It!”. By which he means, “write a program that solves it efficiently”.

Question!

How could knowing about NP-hardness help?

→ Assume further that, after trying in vain for 4 weeks, you got the next meeting with Mr. X. Do you want to say “Um, sorry, but I couldn’t find an efficient solution, please don’t fire me”?

Or would you rather say “Look, I didn’t find an efficient solution. But neither could all the Turing-award winners out there put together, because the problem is NP-hard”?

(This particular **NP** sales pitch is not my invention. This is how Garey and Johnson start their book [Garey and Johnson (1979)].)

Reminder: Algorithmic Problems in Planning

Satisficing Planning

Input: A planning task Π .

Output: A plan for Π , or “unsolvable” if no plan for Π exists.

Optimal Planning

Input: A planning task Π .

Output: An *optimal* plan for Π , or “unsolvable” if no plan for Π exists.

Decision Problems in (STRIPS) Planning

Definition (PlanEx). By *PlanEx*, we denote the problem of deciding, given a STRIPS planning task Π , whether or not there exists a plan for Π .

→ Corresponds to satisficing planning.

Definition (PlanLen). By *PlanLen*, we denote the problem of deciding, given a STRIPS planning task Π and an integer B , whether or not there exists a plan for Π of length at most B .

→ Corresponds to optimal planning.

Definition (PolyPlanLen). By *PolyPlanLen*, we denote the problem of deciding, given a STRIPS planning task Π and an integer B *bounded by a polynomial in the size of Π* , whether or not there exists a plan for Π of length at most B .

→ Corresponds to optimal planning with “small” plans. Example of a planning domain with exponentially long plans? Towers of Hanoi.

Complexity of PlanEx [Bylander (1994)]

Lemma. PlanEx is **PSPACE**-hard.

→ "At least as hard as any other problem contained in **PSPACE**."

Proof Sketch. Given a Turing machine with space bounded by polynomial $p(|w|)$, we can in polynomial time (in the size of the machine) generate an equivalent STRIPS planning task. Say the possible symbols in tape cells are x_1, \dots, x_m and the internal states are s_1, \dots, s_n , accepting state s_{acc} .

- The contents of the tape cells:
 $in(1, x_1), \dots, in(p(|w|), x_1), \dots, in(1, x_m), \dots, in(p(|w|), x_m)$.
- The position of the R/W head: $at(1), \dots, at(p(|w|))$.
- The internal state of the machine: $state(s_1), \dots, state(s_n)$.
- Transitions rules \mapsto STRIPS actions; accepting state \mapsto STRIPS goal $\{state(s_{acc})\}$; initial state obvious.
- This reduction to STRIPS runs in polynomial-time because we need only polynomially many facts.

Complexity of PlanEx, ctd. [Bylander (1994)]

Lemma. *PlanEx is a member of PSPACE.*

→ "At most as hard as any other problem contained in **PSPACE**."

Proof. Because **PSPACE** = **NPSPACE**, it suffices to show that PlanEx is a member of **NPSPACE**:

1. $s := I; l := 0;$
2. Guess an applicable action a , compute the outcome state s' , set $l := l + 1;$
3. If s' contains the goal then succeed;
4. If $l \geq 2^{|P|}$ then fail else goto 2;

→ Remembering the actual action *sequence* would take exponential space in case of exponentially long plans (cf. slide 39). But, to decide PlanEx, we only need to remember its length.

Theorem (Complexity of PlanEx). *PlanEx is PSPACE-complete.*
(Immediate from previous two lemmas)

Complexity of PlanLen [Bylander (1994)]

PlanLen isn't any easier than PlanEx:

Corollary. *PlanLen is PSPACE-complete.*

Proof. Membership: Same as before but failing at $l \geq B$. **Hardness?**
Setting $B := 2^{|P|}$, PlanLen answers PlanEx: If a plan exists, then there exists a plan that traverses each possible state at most once.

PolyPlanLen is easier than PlanEx:

Theorem. *PolyPlanLen is NP-complete.*

Proof. **Membership?** Guess B actions and check whether they form a plan. This runs in polynomial time because B is polynomially bounded.
Hardness: E.g., by reduction from SAT.

→ Bounding plan length does not help in the general case as we can set the bound to a trivial (exponential) upper bound on plan length. If we restrict plan length to be “short” (polynomial), planning becomes easier.


Domain-Specific PlanEx vs. PlanLen ...

... is more interesting than the general case.


- *In general*, both have the same complexity.
- *Within particular applications*, bounded length plan existence is often harder than plan existence.
- This happens in many IPC benchmark domains: PlanLen is **NP**-complete while PlanEx is in **P**.
- For example: Blocksworld and Logistics.

→ PlanEx \approx satisficing planning, PlanLen \approx optimal planning. In practice, optimal planning is (almost) never “easy”.


The Blockworld is Hard?



The Blockworld is Hard!



Initial State




Goal State

Miconic-ADL: PlanEx is Hard



- VIP: Served first.
- D: Lift may only go *down* when inside; similar for U.
- NA: Never-alone; AT: Attendant.
- A, B: Never together in the same elevator (!)
- P: Normal passenger :-)



Summary

- General problem solving attempts to develop solvers that perform well across a large class of problems.
- Planning, as considered here, is a form of general problem solving dedicated to the class of classical search problems. (Actually, we also address inaccessible, stochastic, dynamic, continuous, and multi-agent settings.)
- Heuristic search planning has dominated the International Planning Competition (IPC). We focus on it here.
- STRIPS is the simplest possible, while reasonably expressive, language for our purposes. It uses Boolean variables (facts), and defines actions in terms of precondition, add list, and delete list.
- PDDL is the de-facto standard language for describing planning problems.
- Plan existence (bounded or not) is **PSPACE**-complete to decide for STRIPS. If we bound plans polynomially, we get down to **NP**-completeness.

Reading

- *Chapters 10: Classical Planning and 11: Planning and Acting in the Real World* [Russell and Norvig (2010)].

Content: Although the book is named “A Modern Approach”, the planning section was written long before the IPC was even dreamt of, before PDDL was conceived, and several years before heuristic search hit the scene. As such, what we have right now is the attempt of two outsiders trying in vain to catch up with the dramatic changes in planning since 1995.

Chapter 10 is Ok as a background read. Some issues are, imho, misrepresented, and it's far from being an up-to-date account. But it's Ok to get some additional intuitions in words different from my own.

Chapter 11 is annoyingly named (I've seen lots of classical planning in the “real world”), but is useful in our context here because I don't cover any of it. If you're interested in extended/alternative planning paradigms, do read it.

References I

- Avrim L. Blum and Merrick L. Furst. Fast planning through planning graph analysis. In S. Mellish, editor, *Proceedings of the 14th International Joint Conference on Artificial Intelligence (IJCAI'95)*, pages 1636–1642, Montreal, Canada, August 1995. Morgan Kaufmann.
- Avrim L. Blum and Merrick L. Furst. Fast planning through planning graph analysis. *Artificial Intelligence*, 90(1–2):279–298, 1997.
- Blai Bonet and Héctor Geffner. Planning as heuristic search: New results. In S. Biundo and M. Fox, editors, *Proceedings of the 5th European Conference on Planning (ECP'99)*, pages 60–72. Springer-Verlag, 1999.
- Blai Bonet and Héctor Geffner. Planning as heuristic search. *Artificial Intelligence*, 129(1–2):5–33, 2001.
- Blai Bonet, Lee McCluskey, José Reinaldo Silva, and Brian Williams, editors. *Proceedings of the 22nd International Conference on Automated Planning and Scheduling (ICAPS'12)*. AAAI Press, 2012.
- Tom Bylander. The computational complexity of propositional STRIPS planning. *Artificial Intelligence*, 69(1–2):165–204, 1994.

References II

Carmel Domshlak, Jörg Hoffmann, and Michael Katz. Red-black planning: A new systematic approach to partial delete relaxation. *Artificial Intelligence*, 221:73–114, 2015.

Stefan Edelkamp. Planning with pattern databases. In A. Cesta and D. Borrajo, editors, *Proceedings of the 6th European Conference on Planning (ECP'01)*, pages 13–24. Springer-Verlag, 2001.

Richard E. Fikes and Nils Nilsson. STRIPS: A new approach to the application of theorem proving to problem solving. *Artificial Intelligence*, 2:189–208, 1971.

Maria Fox and Derek Long. PDDL2.1: An extension to PDDL for expressing temporal planning domains. *Journal of Artificial Intelligence Research*, 20:61–124, 2003.

Michael R. Garey and David S. Johnson. *Computers and Intractability—A Guide to the Theory of NP-Completeness*. Freeman, San Francisco, CA, 1979.

Alfonso Gerevini, Alessandro Saetti, and Ivan Serina. Planning through stochastic local search and temporal action graphs. *Journal of Artificial Intelligence Research*, 20:239–290, 2003.

References III

- Alfonso Gerevini, Patrik Haslum, Derek Long, Alessandro Saetti, and Yannis Dimopoulos. Deterministic planning in the fifth international planning competition: PDDL3 and experimental evaluation of the planners. *Artificial Intelligence*, 173(5-6):619–668, 2009.
- Malik Ghallab, Dana Nau, and Paolo Traverso. *Automated Planning: Theory and Practice*. Morgan Kaufmann, 2004.
- Patrik Haslum and Hector Geffner. Admissible heuristics for optimal planning. In S. Chien, R. Kambhampati, and C. Knoblock, editors, *Proceedings of the 5th International Conference on Artificial Intelligence Planning Systems (AIPS'00)*, pages 140–149, Breckenridge, CO, 2000. AAAI Press, Menlo Park.
- Malte Helmert and Carmel Domshlak. Landmarks, critical paths and abstractions: What's the difference anyway? In Alfonso Gerevini, Adele Howe, Amedeo Cesta, and Ioannis Refanidis, editors, *Proceedings of the 19th International Conference on Automated Planning and Scheduling (ICAPS'09)*, pages 162–169. AAAI Press, 2009.

References IV

- Malte Helmert and Hector Geffner. Unifying the causal graph and additive heuristics. In Jussi Rintanen, Bernhard Nebel, J. Christopher Beck, and Eric Hansen, editors, *Proceedings of the 18th International Conference on Automated Planning and Scheduling (ICAPS'08)*, pages 140–147. AAAI Press, 2008.
- Malte Helmert, Patrik Haslum, and Jörg Hoffmann. Flexible abstraction heuristics for optimal sequential planning. In Mark Boddy, Maria Fox, and Sylvie Thiebaux, editors, *Proceedings of the 17th International Conference on Automated Planning and Scheduling (ICAPS'07)*, pages 176–183, Providence, Rhode Island, USA, 2007. Morgan Kaufmann.
- Malte Helmert. The Fast Downward planning system. *Journal of Artificial Intelligence Research*, 26:191–246, 2006.
- Jörg Hoffmann and Stefan Edelkamp. The deterministic part of ipc-4: An overview. *Journal of Artificial Intelligence Research*, 24:519–579, 2005.
- Jörg Hoffmann and Bernhard Nebel. The FF planning system: Fast plan generation through heuristic search. *Journal of Artificial Intelligence Research*, 14:253–302, 2001.

References V

Erez Karpas and Carmel Domshlak. Cost-optimal planning with landmarks. In Craig Boutilier, editor, *Proceedings of the 21st International Joint Conference on Artificial Intelligence (IJCAI'09)*, pages 1728–1733, Pasadena, California, USA, July 2009. Morgan Kaufmann.

Michael Katz, Jörg Hoffmann, and Malte Helmert. How to relax a bisimulation? In Bonet et al. Bonet *et al.* (2012), pages 101–109.

Michael Katz, Jörg Hoffmann, and Carmel Domshlak. Who said we need to relax *all* variables? In Daniel Borrajo, Simone Fratini, Subbarao Kambhampati, and Angelo Oddi, editors, *Proceedings of the 23rd International Conference on Automated Planning and Scheduling (ICAPS'13)*, pages 126–134, Rome, Italy, 2013. AAAI Press.

Henry A. Kautz and Bart Selman. Planning as satisfiability. In B. Neumann, editor, *Proceedings of the 10th European Conference on Artificial Intelligence (ECAI'92)*, pages 359–363, Vienna, Austria, August 1992. Wiley.

Henry A. Kautz and Bart Selman. Pushing the envelope: Planning, propositional logic, and stochastic search. In William J. Clancey and Daniel Weld, editors, *Proceedings of the 13th National Conference of the American Association for Artificial Intelligence (AAAI'96)*, pages 1194–1201, Portland, OR, July 1996. MIT Press.

References VI

- Emil Keyder, Jörg Hoffmann, and Patrik Haslum. Semi-relaxed plan heuristics. In Bonet et al. Bonet *et al.* (2012), pages 128–136.
- Jana Koehler, Bernhard Nebel, Jörg Hoffmann, and Yannis Dimopoulos. Extending planning graphs to an ADL subset. In S. Steel and R. Alami, editors, *Proceedings of the 4th European Conference on Planning (ECP'97)*, pages 273–285. Springer-Verlag, 1997.
- Drew McDermott, Malik Ghallab, Adele Howe, Craig Knoblock, Ashwin Ram, Manuela Veloso, Daniel Weld, and David Wilkins. *The PDDL Planning Domain Definition Language*. The AIPS-98 Planning Competition Comitee, 1998.
- Allen Newell and Herbert Simon. GPS, a program that simulates human thought. In E. Feigenbaum and J. Feldman, editors, *Computers and Thought*, pages 279–293. McGraw-Hill, 1963.
- Raz Nissim, Jörg Hoffmann, and Malte Helmert. Computing perfect heuristics in polynomial time: On bisimulation and merge-and-shrink abstraction in optimal planning. In Toby Walsh, editor, *Proceedings of the 22nd International Joint Conference on Artificial Intelligence (IJCAI'11)*, pages 1983–1990. AAAI Press/IJCAI, 2011.

References VII

- J. Scott Penberthy and Daniel S. Weld. UCPOP: A sound, complete, partial order planner for ADL. In B. Nebel, W. Swartout, and C. Rich, editors, *Principles of Knowledge Representation and Reasoning: Proceedings of the 3rd International Conference (KR-92)*, pages 103–114, Cambridge, MA, October 1992. Morgan Kaufmann.
- Silvia Richter and Matthias Westphal. The LAMA planner: Guiding cost-based anytime planning with landmarks. *Journal of Artificial Intelligence Research*, 39:127–177, 2010.
- Jussi Rintanen, Keijo Heljanko, and Ilkka Niemelä. Planning as satisfiability: parallel plans and algorithms for plan search. *Artificial Intelligence*, 170(12-13):1031–1080, 2006.
- Jussi Rintanen. Heuristics for planning with SAT. In *Proceedings of the 16th International Conference on Principles and Practice of Constraint Programming*, pages 414–428, 2010.
- Stuart Russell and Peter Norvig. *Artificial Intelligence: A Modern Approach (Third Edition)*. Prentice-Hall, Englewood Cliffs, NJ, 2010.