

ARTIFICIAL INTELLIGENCE

CONSTRAINT SATISFACTION PROBLEMS

Summary

- ◇ CSP RN Chapter 5
- ◇ Constraint satisfaction problems
- ◇ Backtracking search for CSP
- ◇ Constraint propagation
- ◇ Local search for CSP
- ◇ Problem Structure

Constraint satisfaction problems (CSPs)

Standard search problem:

state is a “black box” — any data structure that supports goal test, eval, successor

CSP:

State is defined by *variables* X_i with *values* from *domain* D_i

Goal test is a set of *constraints* specifying allowable combinations of values v_i for subsets of variables

Solution is an *assignment* $\{X_i = v_j, \dots\}$ that does not violate the constraints

Key feature: *general-purpose* algorithms with more power than standard search algorithms

Constraint satisfaction problems (CSPs) contd.

A CSP is a triple $\langle X, D, C \rangle$, where:

- $X = \{X_1, \dots, X_n\}$ is the set of variables;
- $D = \{D_1, \dots, D_n\}$ is a set of domains, specifying the allowed values for the variables;
- $C = \{C_1, \dots, C_n\}$ is a set of constraints,
where $C_i = \langle X_i^k, R_i^k \rangle$
 X_i^k is a subset of k elements of X
 R_i^k is a k -order relation over X_i^k

Varieties of CSPs

Discrete variables

- ◇ **finite domains**, are the most common
size $d \Rightarrow O(d^n)$ complete assignments
e.g., **Boolean** CSPs, (i.e., variables in $\{true, false\}$)
includes Boolean satisfiability (NP-complete)

Other kinds of CSPs

◇ infinite domains (integers, strings, etc.)

CSPs cannot be solved by enumerating assignments

e.g., job scheduling, var = start/end days for each job

need a constraint language,

e.g., $StartJob_1 + 5 \leq StartJob_3$

linear constraints solvable, **nonlinear** undecidable

Infinite domains can be reduced to finite by putting an upper bound to values.

Other kinds of of CSPs contd.

◇ Continuous variables

e.g., start/end times for Hubble Telescope observations
linear constraints solvable in poly time by
Linear Programming methods

Continuous domains are not considered here

⇒ **Operation Research**

Ariety of constraints

Unary constraints involve a single variable,

e.g., $SA \neq green$

Binary constraints involve pairs of variables,

e.g., $SA \neq WA$

Higher-order constraints involve 3 or more variables,

e.g. $alldiff(X_1, \dots, X_n)$

Higher-order constraints can be reduced to binary constraint (by increasing the number of variables and of constraints)

Unary constraints can be treated as domain restrictions

Preference constraints

Preferences or soft constraints (versus **absolute** constraints),
e.g., *red* is better than *green*,
often representable by a cost for each variable assignment
→ constrained optimization problems (COP)

Example: Map-Coloring



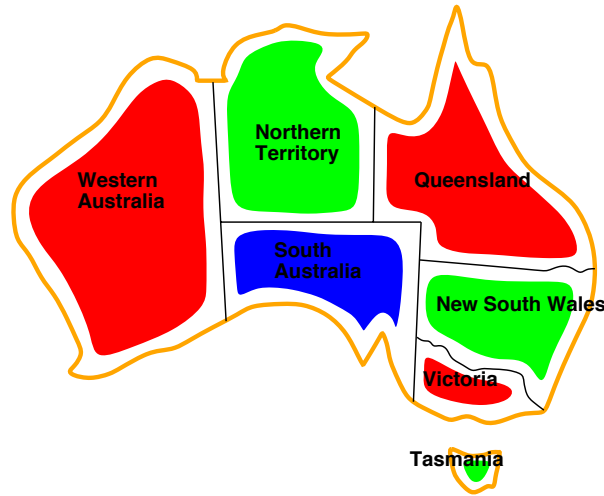
Variables WA, NT, Q, NSW, V, SA, T

Domains $D_i = \{red, green, blue\}$

Constraints: adjacent regions must have different colors

e.g., $WA \neq NT$ (if the language allows this), or
 $(WA, NT) \in \{(red, green), (red, blue), (green, red), (green, blue), \dots\}$

Example: Map-Coloring contd.

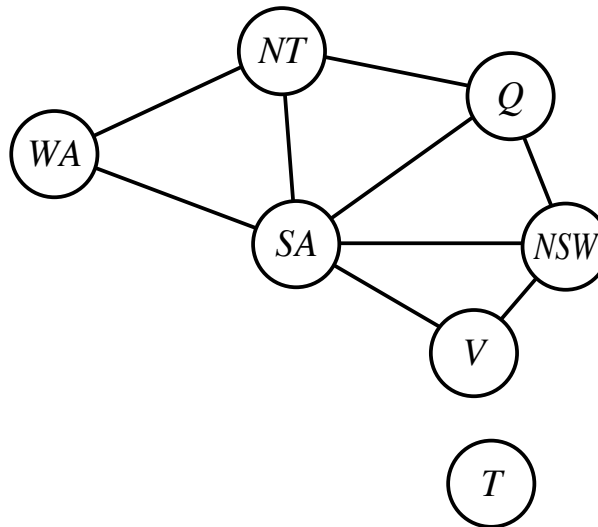


Solutions are assignments satisfying all constraints, e.g.,
 $\{WA = red, NT = green, Q = red, NSW = green, V = red, SA = blue, T = green\}$

Constraint graph

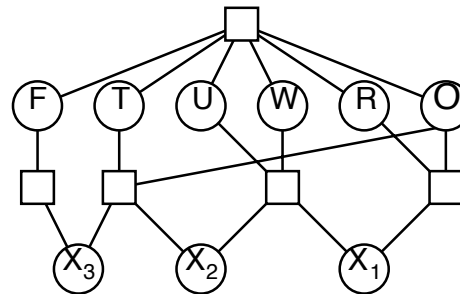
Binary CSP: each constraint relates at most two variables

Constraint graph: nodes are variables, arcs show constraints



General-purpose CSP algorithms use the graph structure to speed up search. E.g., Tasmania is an independent subproblem!

Example: Cryptarithmic

$$\begin{array}{r} \text{ T W O} \\ + \text{ T W O} \\ \hline \text{ F O U R} \end{array}$$


Variables: $F \ T \ U \ W \ R \ O \ X_1 \ X_2 \ X_3$

Domains: $\{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$

Constraints

$alldiff(F, T, U, W, R, O)$

$O + O = R + 10 \cdot X_1, \dots \text{ etc.}$

Example: Sudoku

	1	2	3	4	5	6	7	8	9
A			3		2		6		
B	9			3		5			1
C			1	8		6	4		
D			8	1		2	9		
E	7								8
F			6	7		8	2		
G			2	6		9	5		
H	8			2		3			9
I			5		1		3		

(a)

	1	2	3	4	5	6	7	8	9
A	4	8	3	9	2	1	6	5	7
B	9	6	7	3	4	5	8	2	1
C	2	5	1	8	7	6	4	9	3
D	5	4	8	1	3	2	9	7	6
E	7	2	9	5	6	4	1	3	8
F	1	3	6	7	9	8	2	4	5
G	3	7	2	6	8	9	5	1	4
H	8	1	4	2	5	3	7	6	9
I	6	9	5	4	1	7	3	8	2

(b)

Variables: $A_1 \dots A_9, \dots, I_1 \dots I_9$

Domains: $\{1, 2, 3, 4, 5, 6, 7, 8, 9\}$

Constraints

$alldiff(A_1 \dots A_9), \dots$ etc. $alldiff(A_1 \dots I_1), \dots$ etc.

$alldiff(A_1 \dots A_3, B_1 \dots B_3, C_1 \dots C_3), \dots$ etc.

Real-world CSPs

Assignment problems (e.g., who teaches what class)

Timetabling problems (e.g., which class when and where)

Hardware configuration

Vehicle routing

Transportation scheduling

Factory scheduling

Floorplanning

Several real-world problems involve real-valued variables

Standard search formulation (incremental)

Straightforward approach:

States are defined by the values assigned so far

- ◇ **Initial state**: the empty assignment, $\{ \}$
- ◇ **Successor function**: assign value to unassigned variable that does not conflict with current assignment.
 \Rightarrow fail if no legal assignments (not fixable!)
- ◇ **Goal test**: the current assignment is complete

Standard search formulation consequences

- 1) This is the same for all CSPs!
- 2) Solutions are all at depth n with n variables \Rightarrow use DFS
- 3) Path is irrelevant, so can also use complete-state formulation
- 4) branching $b = (n - \ell)d$ at depth ℓ , hence $n!d^n$ leaves!!!!

Backtracking search

Variable assignments are **commutative**, i.e.,

$[WA = red \text{ then } NT = green]$ same as
 $[NT = green \text{ then } WA = red]$

Only need to consider assignments to a single variable at each node

$\Rightarrow b = d$ and there are d^n leaves

Depth-first search for CSPs with single-variable assignments is called **backtracking** search

Backtracking search is the basic uninformed algorithm for CSPs

Can solve n -queens for $n \approx 25$

Backtracking search contd.

Backtracking = depth search

- 1) fixed variable order
- 2) only legal successors

Backtracking search

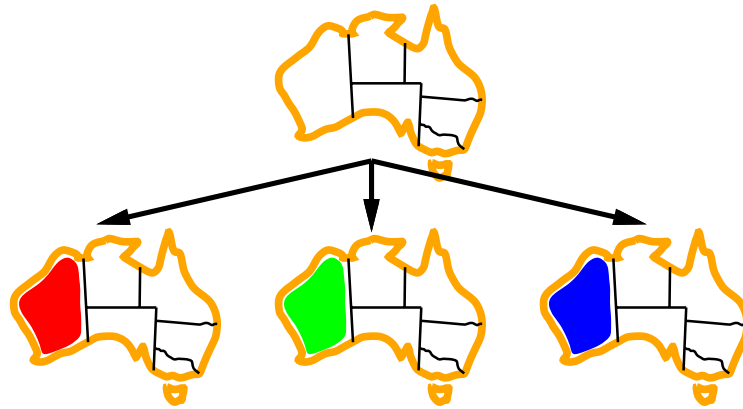
```
function BACKTRACKING-SEARCH(csp) returns solution/failure
  return RECURSIVE-BACKTRACKING( $\{ \}$ , csp)

function RECURSIVE-BACKTRACKING(assignment, csp) returns soln/failure
  if assignment is complete then return assignment
  var  $\leftarrow$  SELECT-UNASSIGNED-VARIABLE(VARIABLES[csp], assignment, csp)
  for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
    if value is consistent with assignment given CONSTRAINTS[csp] then
      add  $\{var = value\}$  to assignment
      result  $\leftarrow$  RECURSIVE-BACKTRACKING(assignment, csp)
      if result  $\neq$  failure then return result
      remove  $\{var = value\}$  from assignment
  return failure
```


Backtracking example




Backtracking example



Backtracking example



Backtracking example



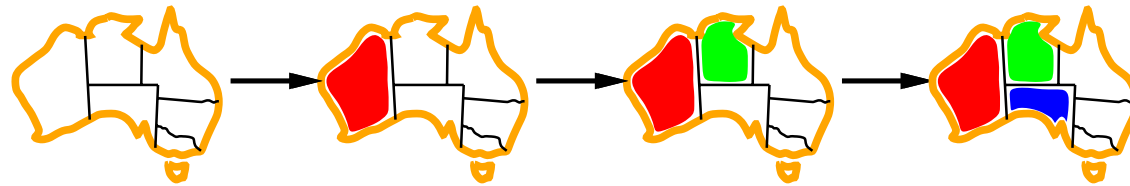
Improving backtracking efficiency

General-purpose methods can give huge gains in speed:

1. Which variable should be assigned next?
 2. In what order should its values be tried?
 3. Can we detect inevitable failure early?
 4. Can we save or reuse partial results of the search?
- ◇ Can we take advantage of problem structure?

Choosing the variable: MRV

Minimum remaining values (MRV):
choose the variable with the fewest legal values

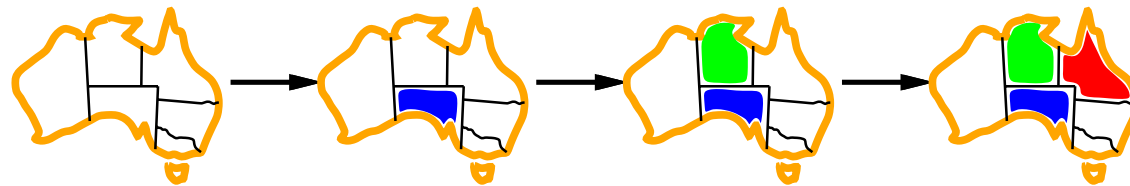


Choosing the variable: Degree heuristic

Tie-breaker among MRV variables

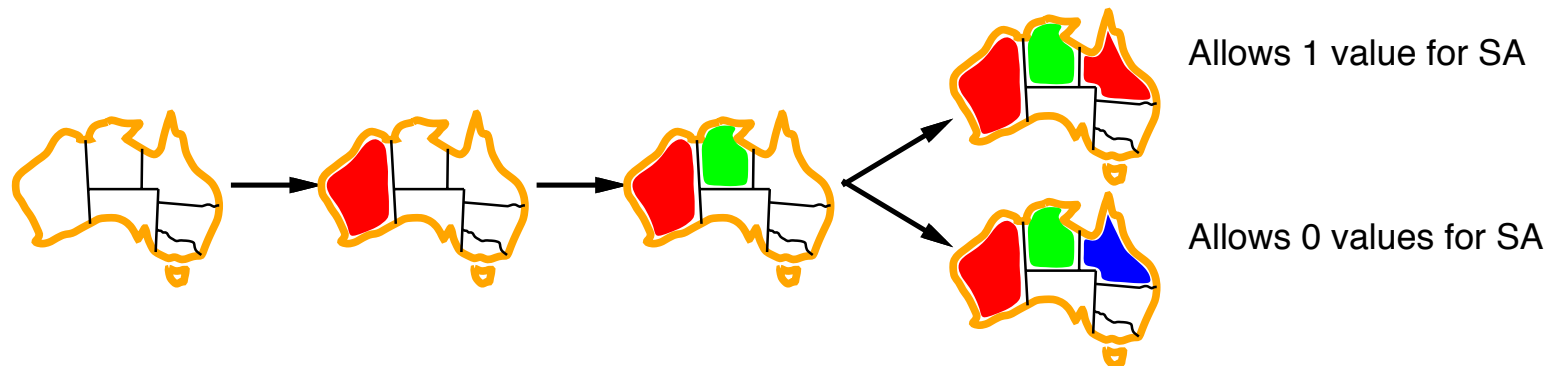
Degree heuristic:

choose the variable with the most constraints on remaining variables



Choosing the value: Least constraining value

Given a variable, choose the least constraining value:
the one that rules out the fewest values in the remaining variables



Combining these heuristics makes 1000 queens feasible

Improving backtracking

◇ Intelligent backtracking (vs chronological),

e.g., conflict directed backjumping

Keep track of the sets of conflicting variables

Forward checking

Idea: Keep track of remaining legal values for unassigned variables

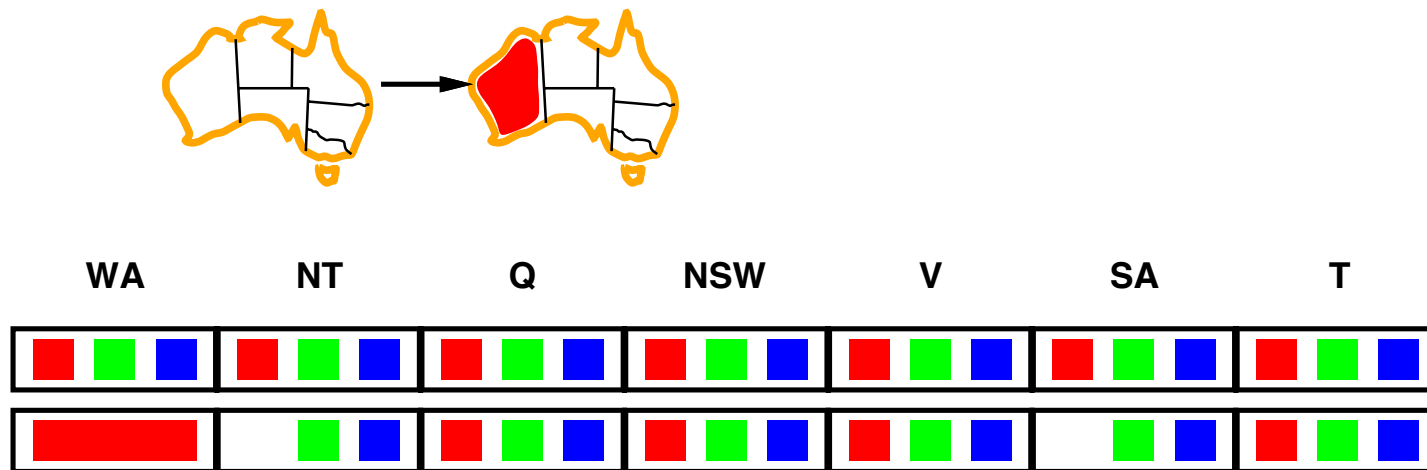
Terminate search when any variable has no legal values



Forward checking

Idea: Keep track of remaining legal values for unassigned variables

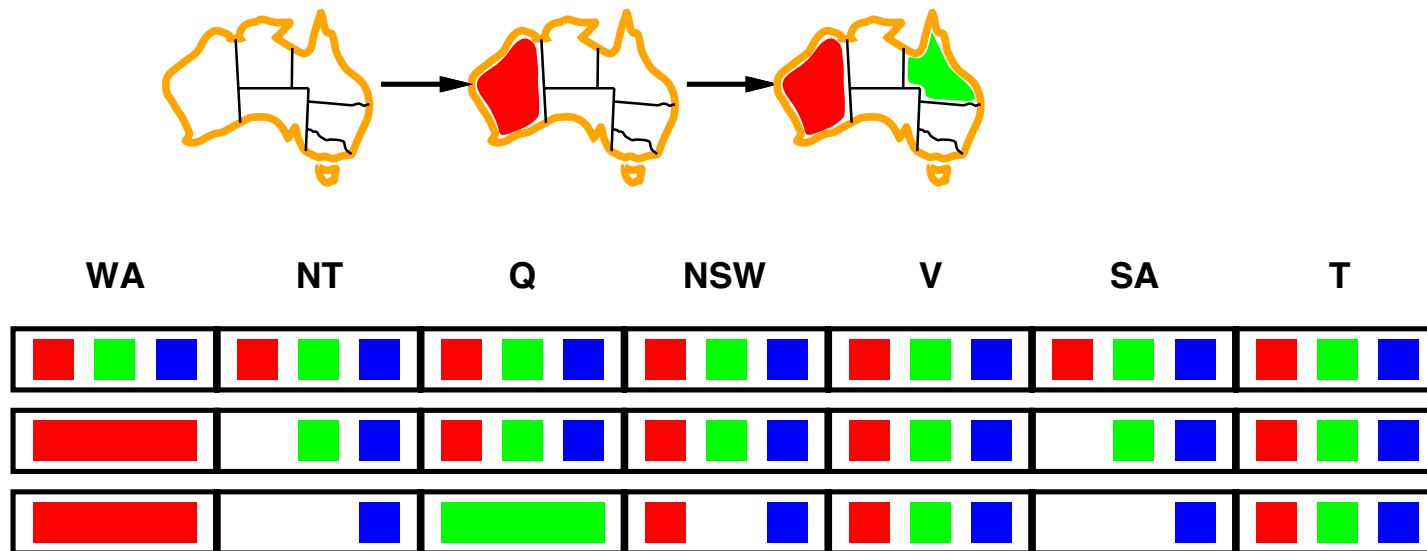
Terminate search when any variable has no legal values



Forward checking

Idea: Keep track of remaining legal values for unassigned variables

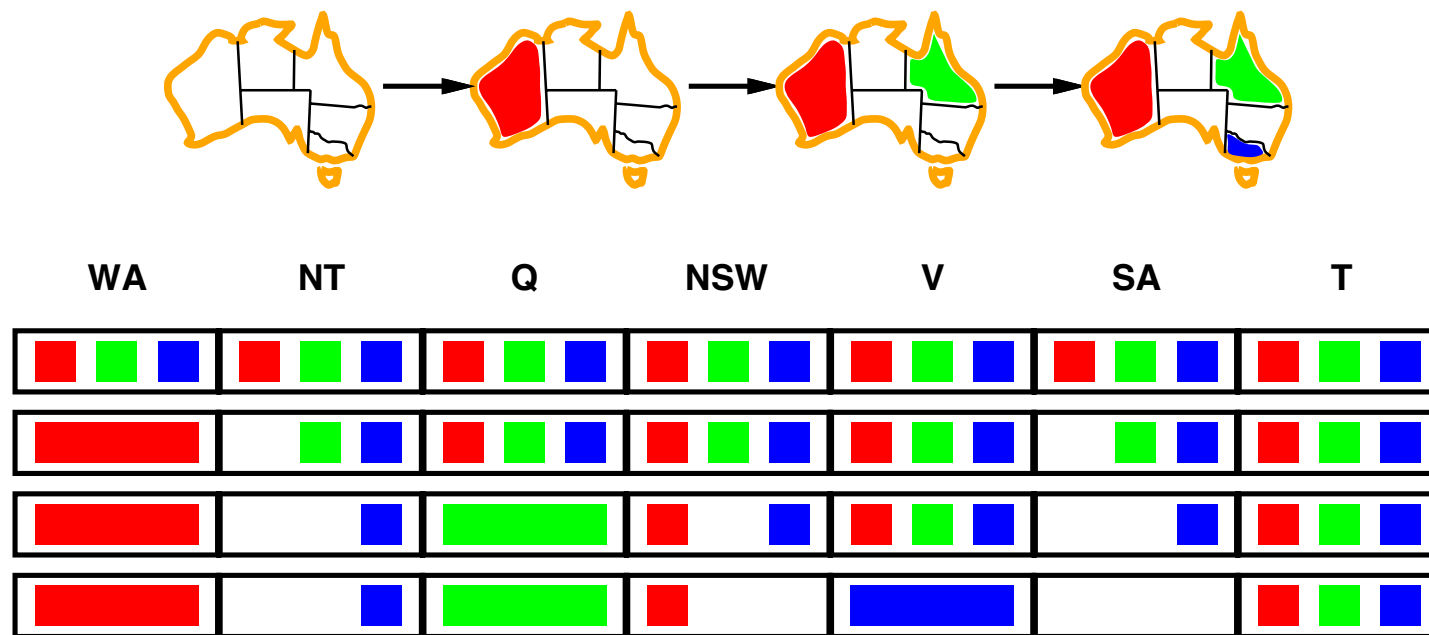
Terminate search when any variable has no legal values



Forward checking


Idea: Keep track of remaining legal values for unassigned variables

Terminate search when any variable has no legal values



Constraint propagation

Forward checking propagates information from assigned to unassigned variables, but doesn't provide early detection for all failures:




NT and SA cannot both be blue!

Constraint propagation repeatedly enforces constraints locally

Arc consistency

Simplest form of propagation makes each arc **consistent**


$X \rightarrow Y$ is consistent iff
for **every** value x of X there is **some** allowed y



Arc consistency

Simplest form of propagation makes each arc **consistent**

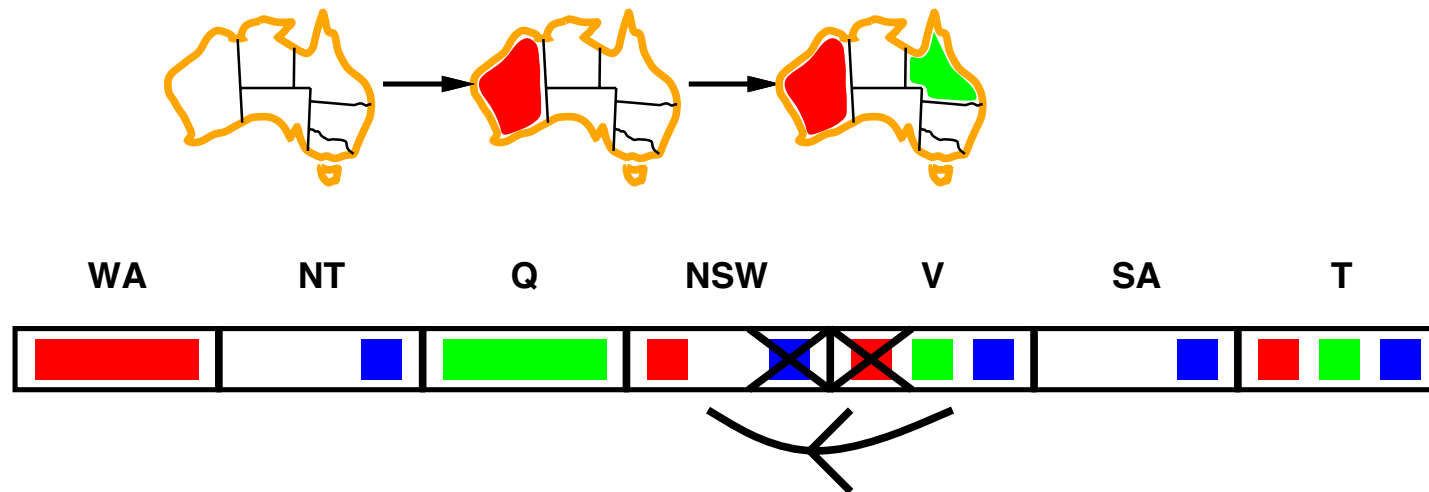
$X \rightarrow Y$ is consistent iff
for **every** value x of X there is **some** allowed y



Arc consistency

Simplest form of propagation makes each arc **consistent**

$X \rightarrow Y$ is consistent iff
for **every** value x of X there is **some** allowed y



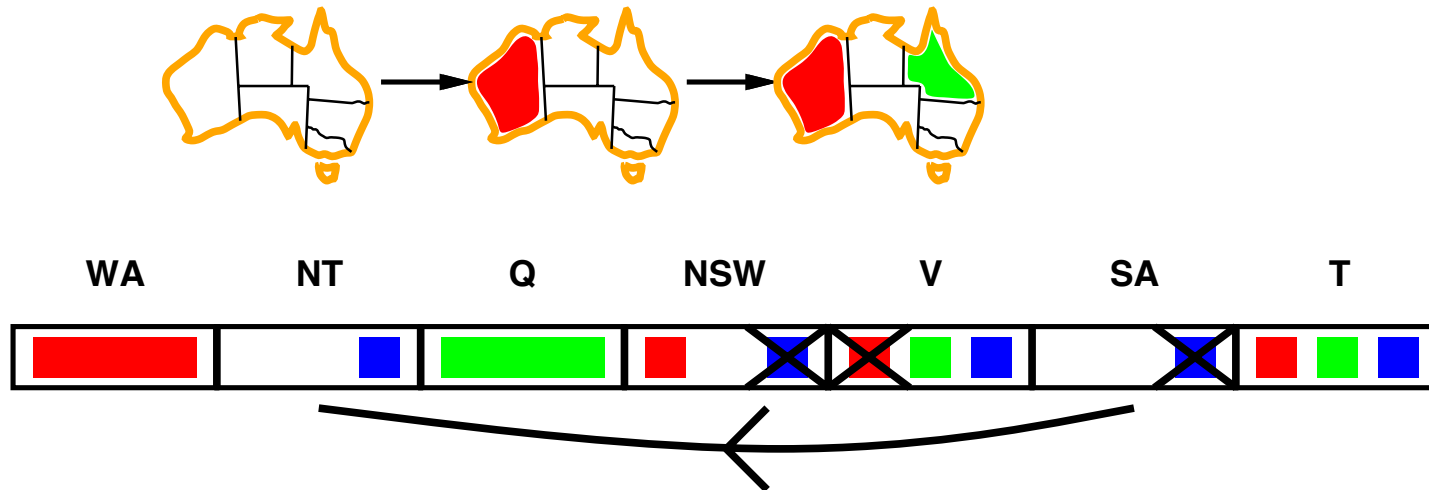
If X loses a value, neighbors of X need to be rechecked

Arc consistency

Simplest form of propagation makes each arc consistent

$X \rightarrow Y$ is consistent iff

for **every** value x of X there is **some** allowed y



If X loses a value, neighbors of X need to be rechecked

Arc consistency detects failure earlier than forward checking

Arc consistency algorithm

function AC-3(*csp*) **returns** false if an inconsistency is found and true otherwise

queue \leftarrow a queue of arcs, initially all the arcs in *csp*

while *queue* is not empty **do**

 (X_i, X_j) \leftarrow POP(*queue*)

if REVISE(*csp*, X_i, X_j) **then**

if size of $D_i = 0$ **then return** *false*

for each X_k **in** X_i .NEIGHBORS - $\{X_j\}$ **do**

 add (X_k, X_i) to *queue*

return *true*

function REVISE(*csp*, X_i, X_j) **returns** true iff we revise the domain of X_i

revised \leftarrow *false*

for each x **in** D_i **do**

if no value y in D_j allows (x, y) to satisfy the constraint between X_i and X_j **then**

 delete x from D_i

revised \leftarrow *true*

return *revised*

Arc consistency

Arc consistency:

- $O(cd^3)$ (with c number of constraints)
- Can be run as a preprocessor or
- after each assignment: in Recursive-Backtracking after adding a new value an **inference** step is performed
 - **forward checking** (removing values from the nodes constrained by the assigned var)
 - **arc consistency** (starting with the arcs that connect the assigned variable with unassigned ones)

Including inference in backtracking search

function BACKTRACKING-SEARCH(*csp*) **returns** a solution or *failure*
 return BACKTRACK(*csp*, { })

function BACKTRACK(*csp*, *assignment*) **returns** a solution or *failure*
 if *assignment* is complete **then return** *assignment*
 var \leftarrow SELECT-UNASSIGNED-VARIABLE(*csp*, *assignment*)
 for each *value* **in** ORDER-DOMAIN-VALUES(*csp*, *var*, *assignment*) **do**
 if *value* is consistent with *assignment* **then**
 add {*var* = *value*} to *assignment*
 inferences \leftarrow INFERENCE(*csp*, *var*, *assignment*)
 if *inferences* \neq *failure* **then**
 add *inferences* to *csp*
 result \leftarrow BACKTRACK(*csp*, *assignment*)
 if *result* \neq *failure* **then return** *result*
 remove *inferences* from *csp*
 remove {*var* = *value*} from *assignment*
 return *failure*

Generalizing

- ◇ path-consistency (3-consistency)
- ◇ k -consistency

Constraints that involve k variables.

with $k = n$ it is possible to obtain a complete inference!!

but computational complexity becomes exponential
(as the problem has an exponential complexity—boolean csp)

Local search algorithms for CSPs

Hill-climbing, simulated annealing typically work with “complete” states, i.e., all variables assigned

To apply to CSPs:

- allow states with unsatisfied constraints

- operators *reassign* variable values

Variable selection: randomly select any conflicted variable

Value selection by *min-conflicts* heuristic:

- choose value that violates the fewest constraints

- i.e., hillclimb with $h(n) =$ total number of violated constraints




Example: 8-Queens

Variables: Y_1, \dots, Y_8

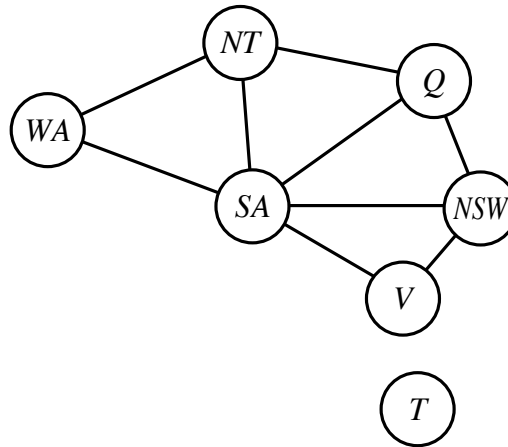
Domains: $\{1, 2, 3, 4, 5, 6, 7, 8\}$

Constraints

$alldiff(Y_1, \dots, Y_8), \dots$



Problem structure



Tasmania and mainland are **independent subproblems**

Identifiable as **connected components** of constraint graph

Problem structure contd.

Suppose each subproblem has c variables out of n total


Worst-case solution cost is $n/c \cdot d^c$, **linear** in n

E.g., $n = 80$, $d = 2$, $c = 20$

$2^{80} = 4$ billion years at 10 million nodes/sec

$4 \cdot 2^{20} = 0.4$ seconds at 10 million nodes/sec

Tree-structured CSPs



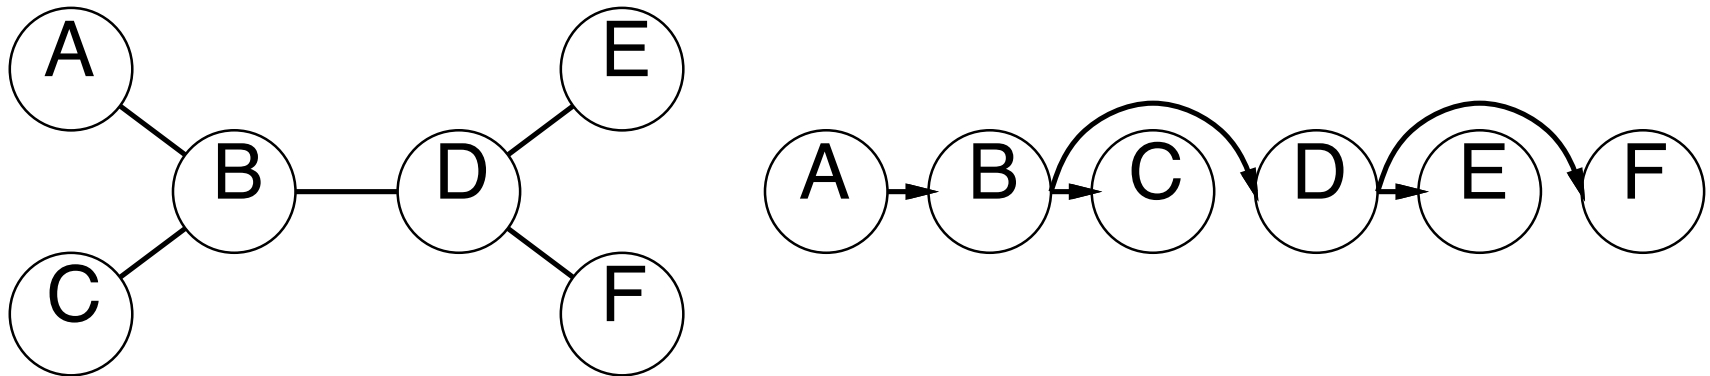
If the constraint graph is a tree, the CSP can be solved in $O(nd^2)$ time

Compare to general CSPs, where worst-case time is $O(d^n)$

This property also applies to logical and probabilistic reasoning: an important example of the relation between syntactic restrictions and the complexity of reasoning.

Algorithm for tree-structured CSPs

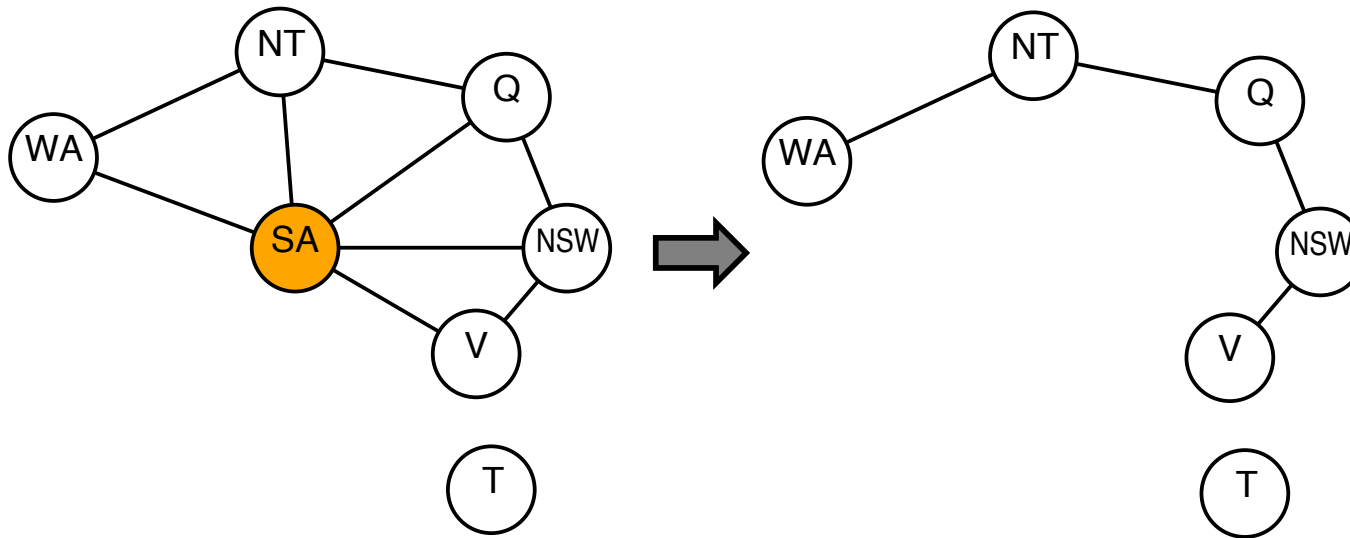
1. Choose a variable as root, order variables from root to leaves such that every node's parent precedes its children nodes



2. For j from n down to 2 , apply $\text{ARCCONSISTENT}(\text{Parent}(X_j), X_j)$
3. For j from 1 to n , assign X_j consistently with $\text{Parent}(X_j)$

Nearly tree-structured CSPs

Conditioning: instantiate a variable, prune its neighbors' domains



Cutset conditioning: instantiate (in all ways) a set of variables such that the remaining constraint graph is a tree

More on problem structure

- Tree decomposition (transformation of graphs into trees)
- Value symmetry (breaking the symmetry to reduce the number of choices)

All problem transformations lead to a polynomial method to solve the problem, but the **exponential** remains in finding the optimal transformation!

Summary

CSPs are a special kind of problem:

- states defined by values of a fixed set of variables

- goal test defined by **constraints** on variable values

Backtracking = depth-first search with one variable assigned per node

Variable ordering and value selection heuristics help significantly

Forward checking prevents assignments that guarantee later failure

Constraint propagation (e.g., arc consistency) does additional work to constrain values and detect inconsistencies

The CSP representation allows analysis of problem structure